AOC Proposal
Systolic Tensor Array

Group｜把家齊高中還給家齊

Member

電機所114　張育誠 N26121892

工科系113　劉子齊 E94096209

電機系113　洪翊碩 E24091027

電機系113　陳喬雅 E14093140

電機系113　鄭喆嚴 E14096724

TABLE OF CONTENTS

01

Model, Quantization and Pruning

Model

VGG16:

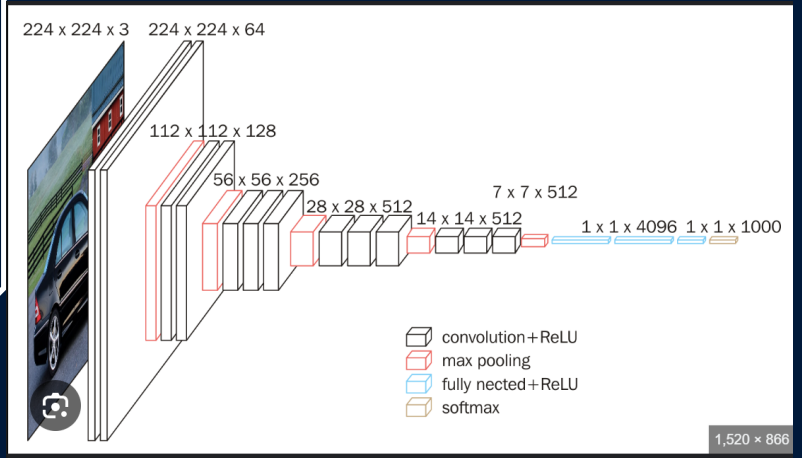

Mobile NetV3:

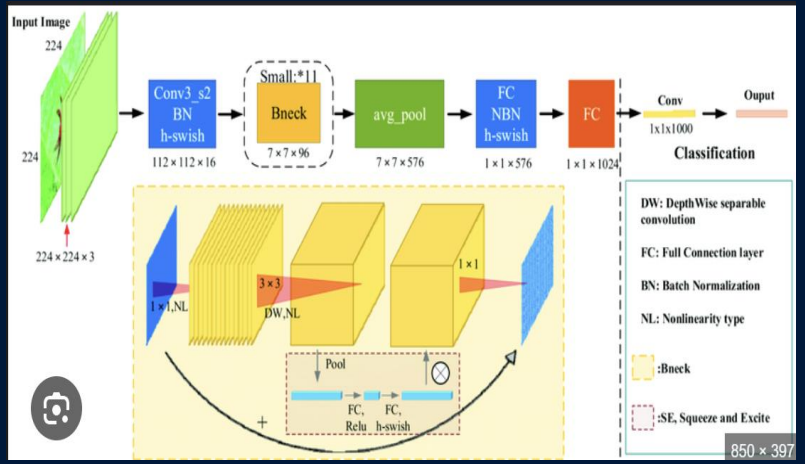

Quantization

QAT:

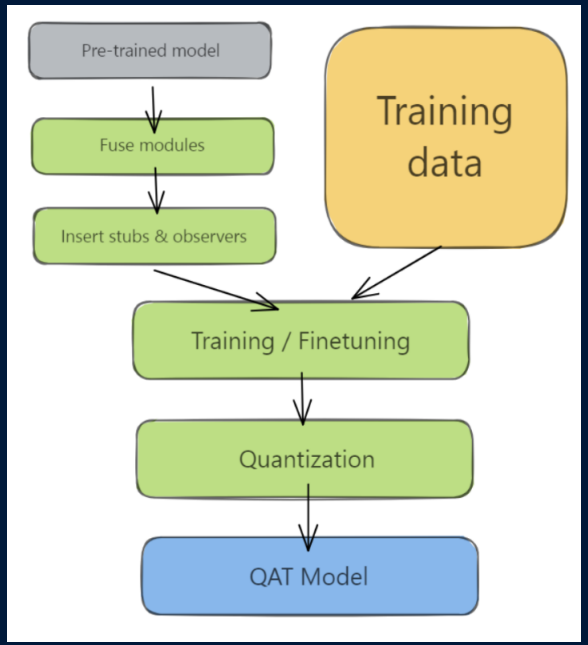

PTQ:

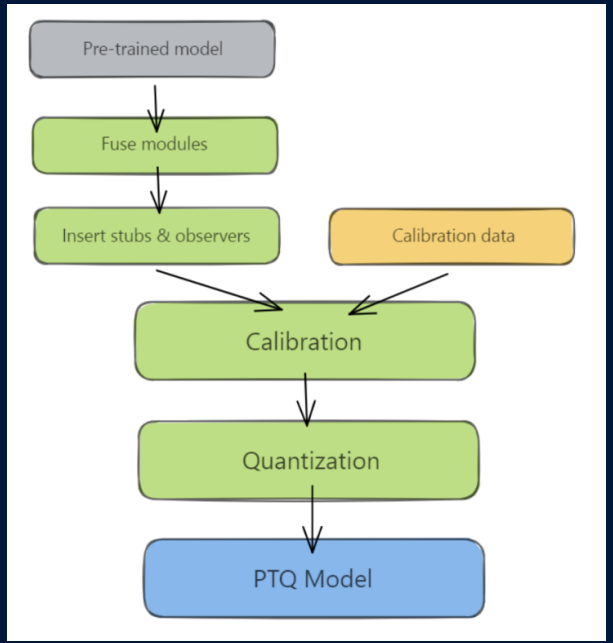

Quantization

QAT:

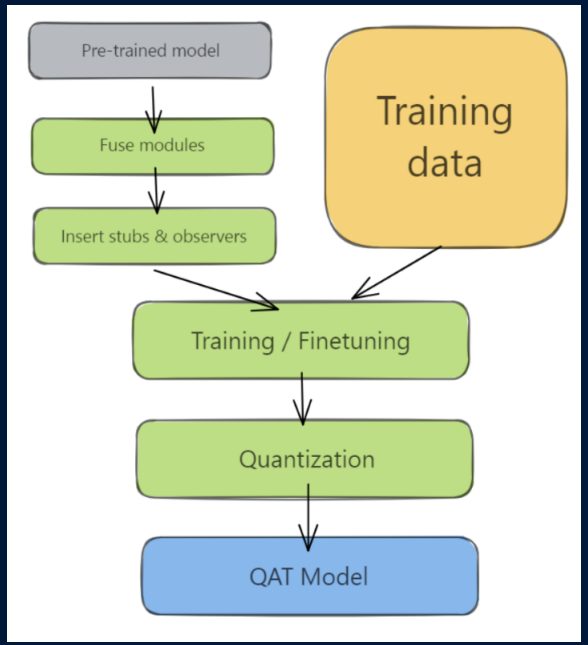

PTQ:

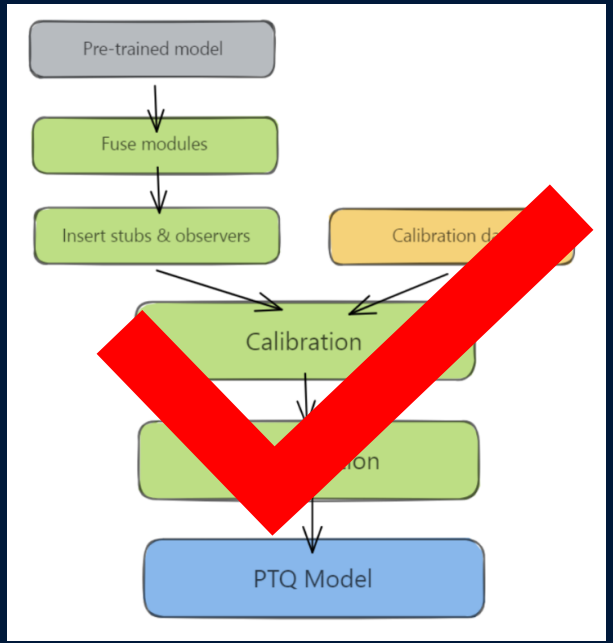

Pruning

Fix position pruning:

```
if  SPARSITY:
        #  updateConv2D()
        #  train(1)
        test(sparsity_model)
        maskConv2D()
        test(model)
        test(sparsity_model)
```

Pruning

Original

```
        [[-1.2915e-04, -9.6068e-04, -1.1363e-03],
         [ 2.1713e-02,  7.5816e-03,  5.7231e-03],
         [ 2.7600e-03,  2.2870e-02,  3.2581e-03]],

        [[ 1.3669e-04,  3.5387e-02,  3.6681e-04],
         [-8.6679e-04,  3.6982e-03,  3.2041e-02],
         [ 5.4913e-03, -5.7838e-03,  1.4623e-03]],

        [[-9.4173e-04,  2.0093e-04,  6.6950e-04],
         [ 4.0101e-03, -1.5172e-02, -9.6157e-03],
         [ 5.9197e-03, -1.0201e-03, -1.6133e-02]]]], device='cuda:0')

Test set: Average loss: 0.3102, Accuracy: 9279/10000 (92.8%)


Test set: Average loss: 0.3081, Accuracy: 9283/10000 (92.8%)
```

Pruning

Saliency-based pruning:

```
if  SPARSITY:
        updateConv2D()
        train(1)
        test(sparsity_model)
        maskConv2D()
        test(model)
        test(sparsity_model)
```

Pruning

Saliency-based pruning:

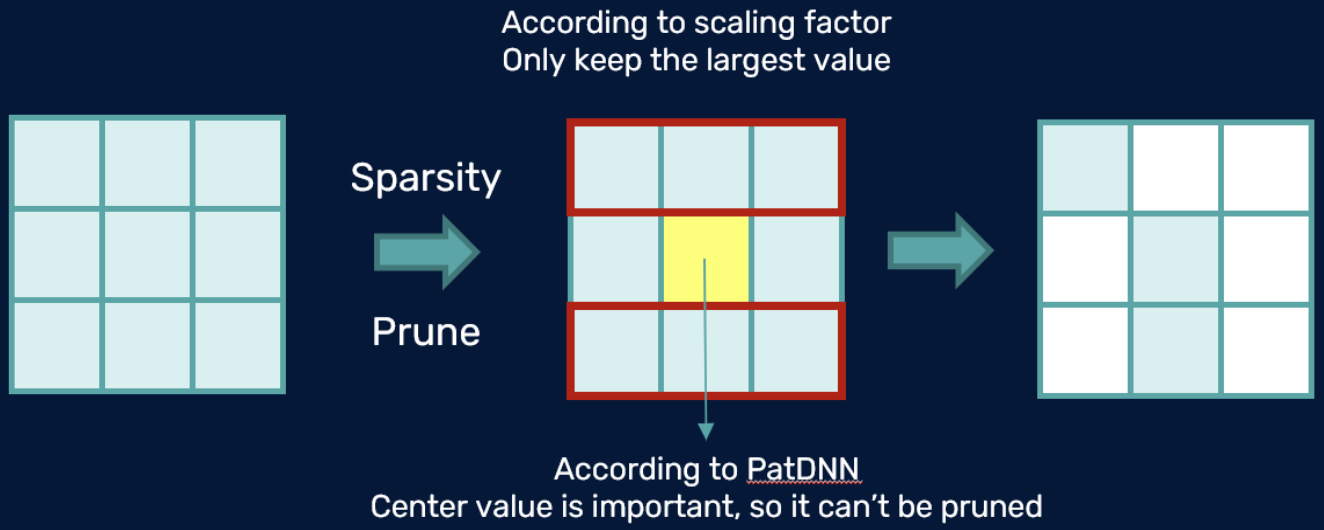

02

PE Architecture and Dataflow

GEMM Accelerator

For Dense Matrix

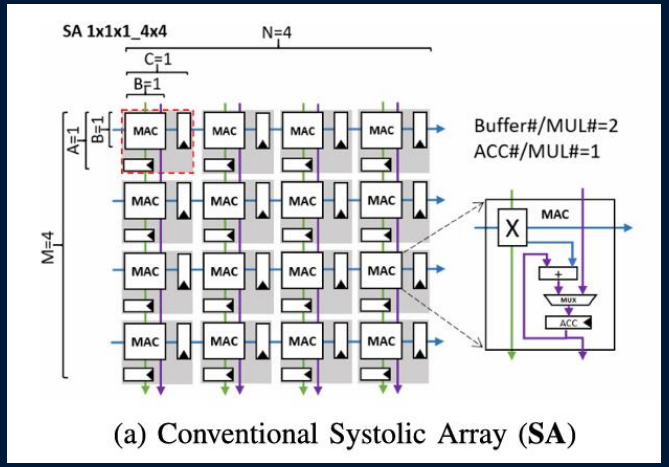

(a) Conventional Systolic Array (**SA**)

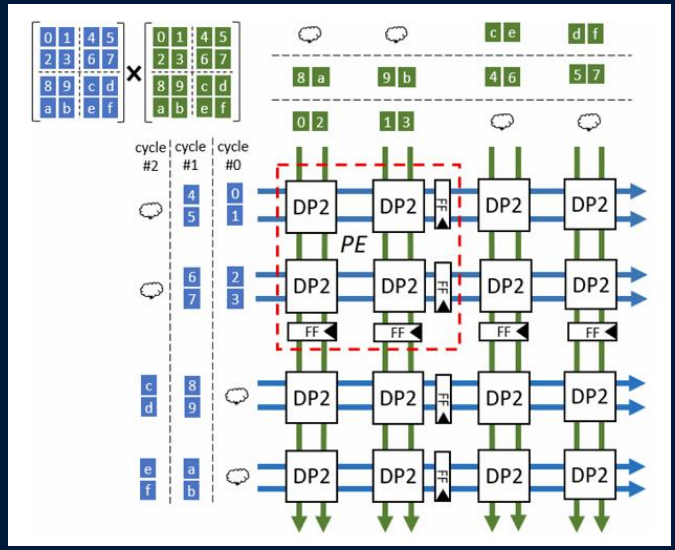

Systolic Tensor Array

GEMM Accelerator

For Sparse Matrix

DBB | Density-Bound Block

Constraint mechanism for
the number of non-zero values
permissible in each block of a matrix.

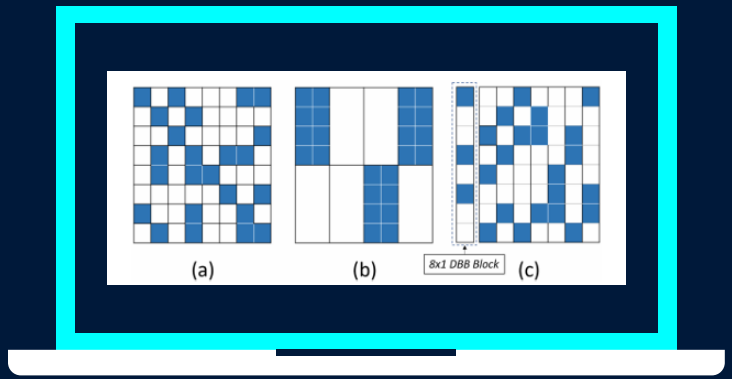

GEMM Accelerator

For Sparse Matrix

Our DBB Method

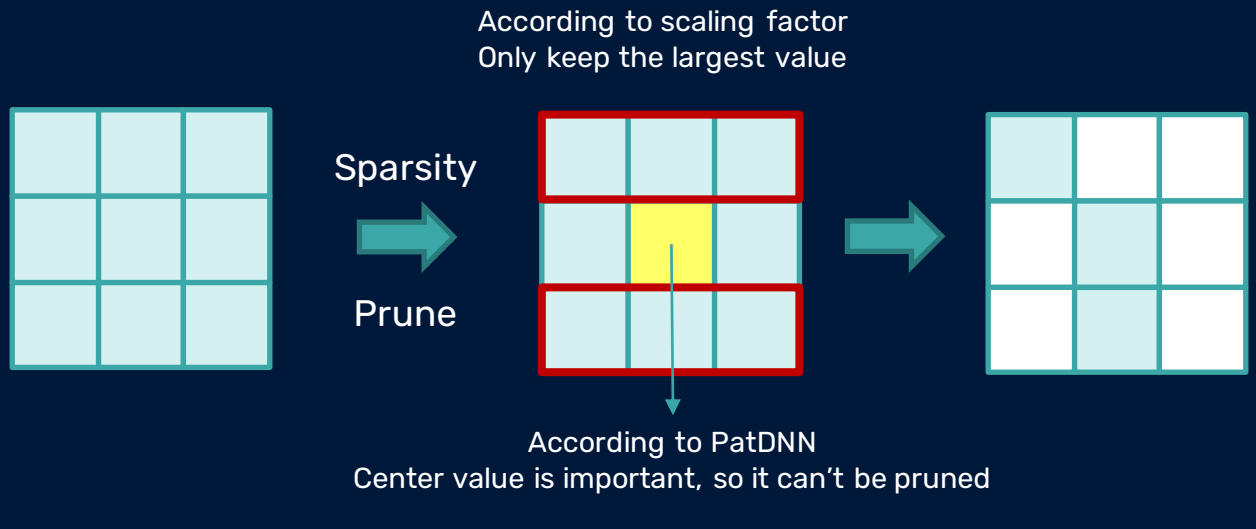

GEMM Accelerator

For Sparse Matrix

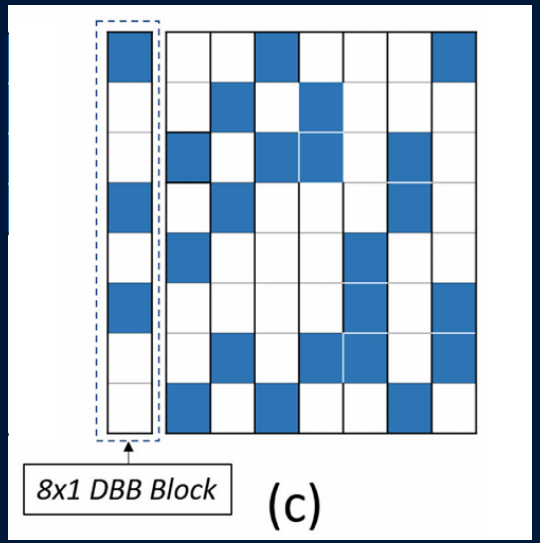

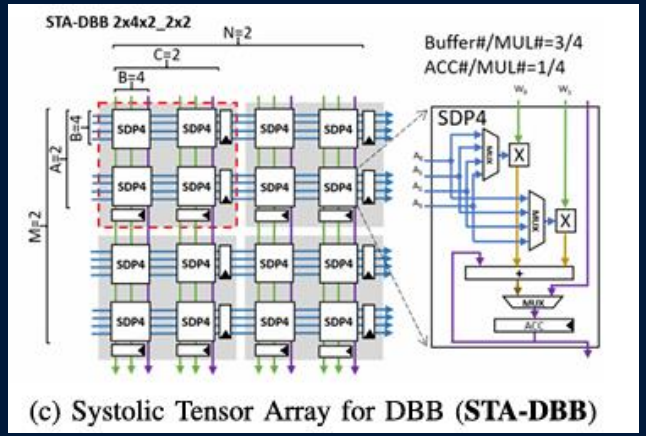

PE Architecture

I_0	I_1	I_2
I_3	I_4	I_5
I_6	I_7	I_8

W_0	W_0	W_0
	W_1	
W_2	W_2	W_2

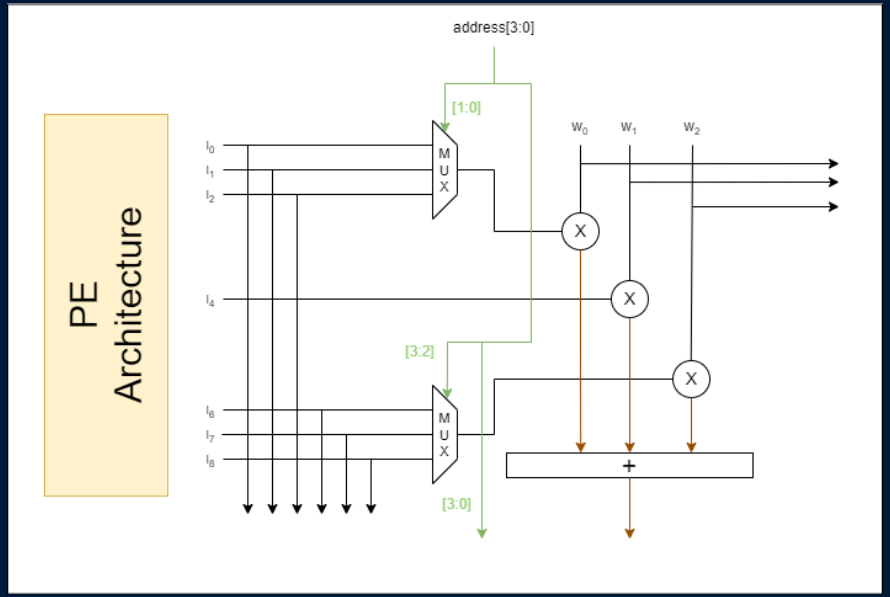

PE I/O Ports

Signal	I/O	Bits	Function
ifmap_i	Input	56 bits	I_0 -> [7:0] I_0, I_1, I_2, I_4, I_6, I_7, I_8
weight_i	Input	24 bits	W_0 -> [7:0] W_0, W_1, W_2
address_i	Input	4 bits	[1:0] for upper weight [3:2] for lower weight
ifmap_o	Output	56 bits	
weight_o	Output	24 bits	
address_o	Output	4 bits	
result	Output	18 bits	Product and Add

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

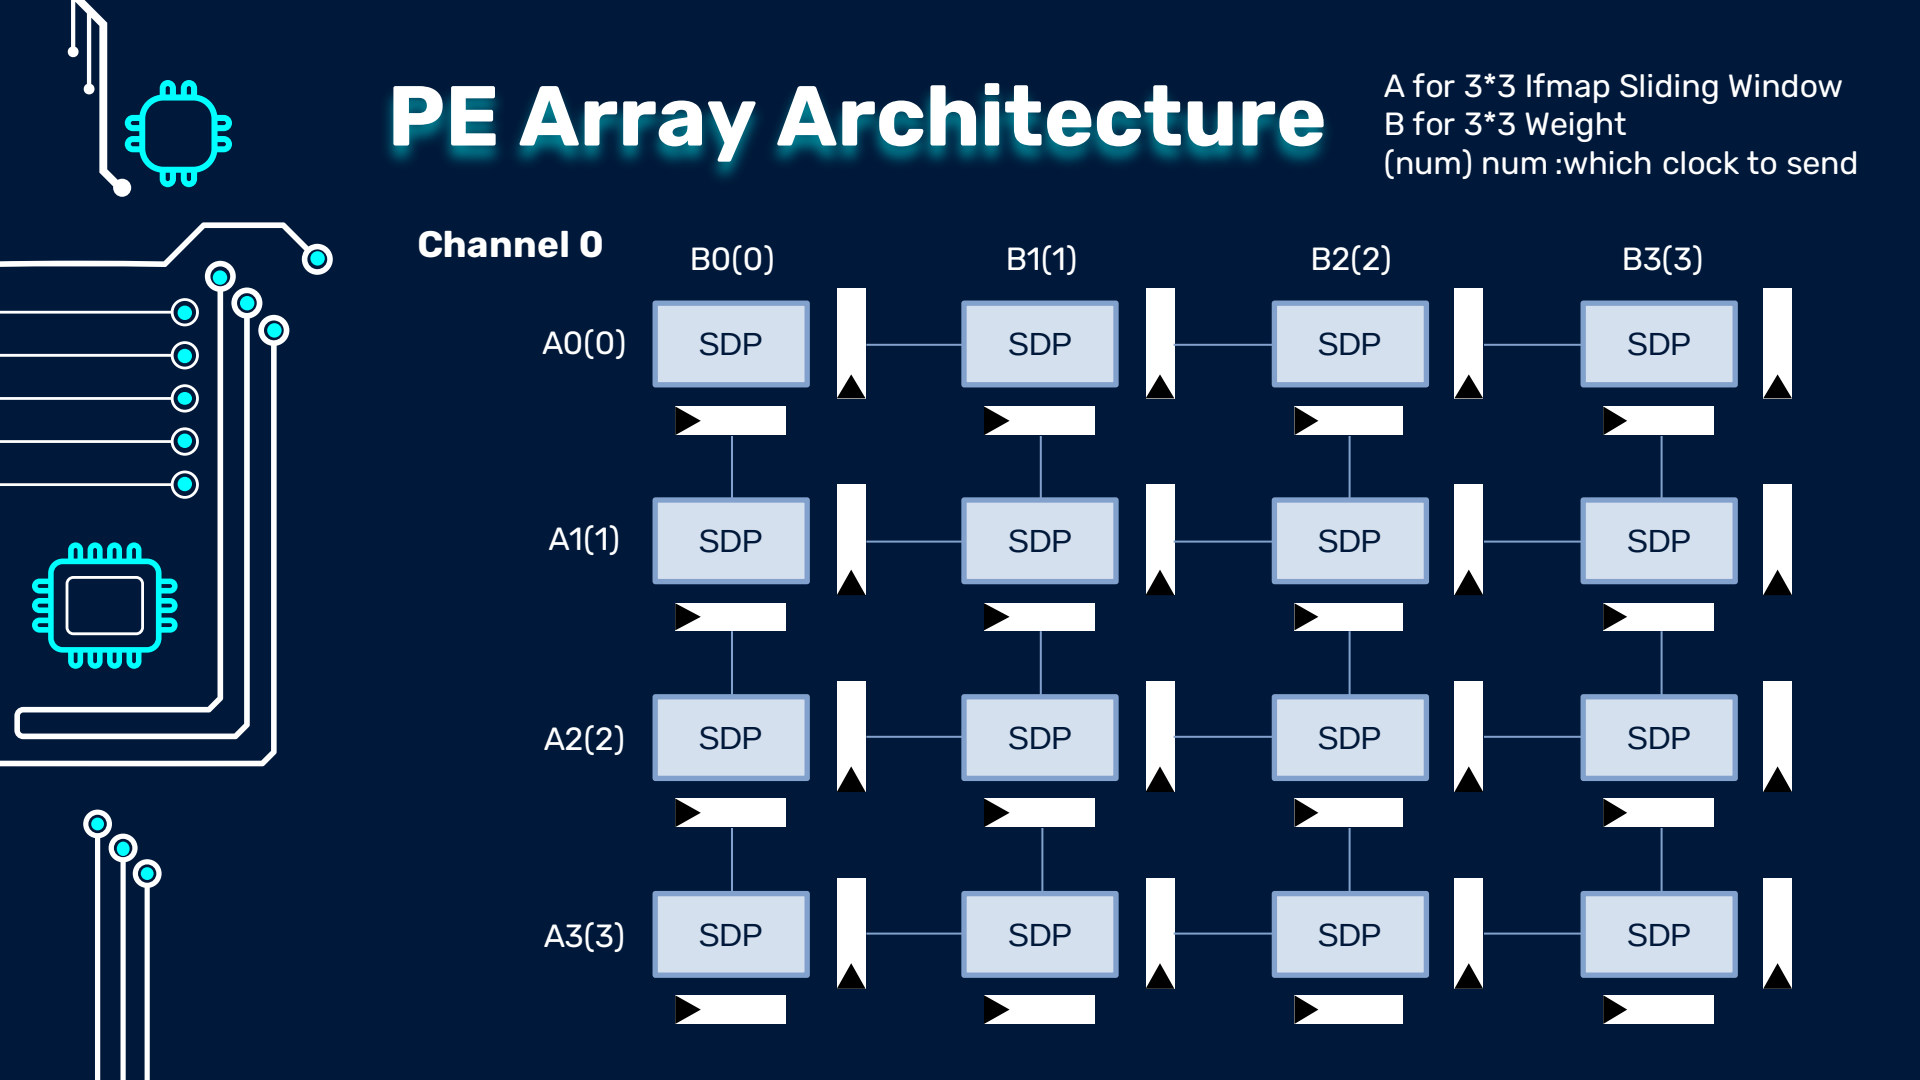

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

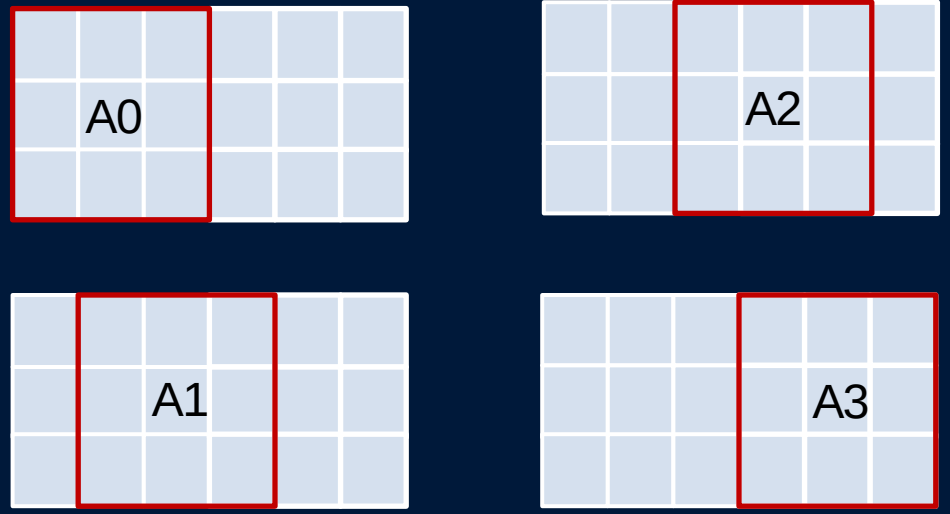

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

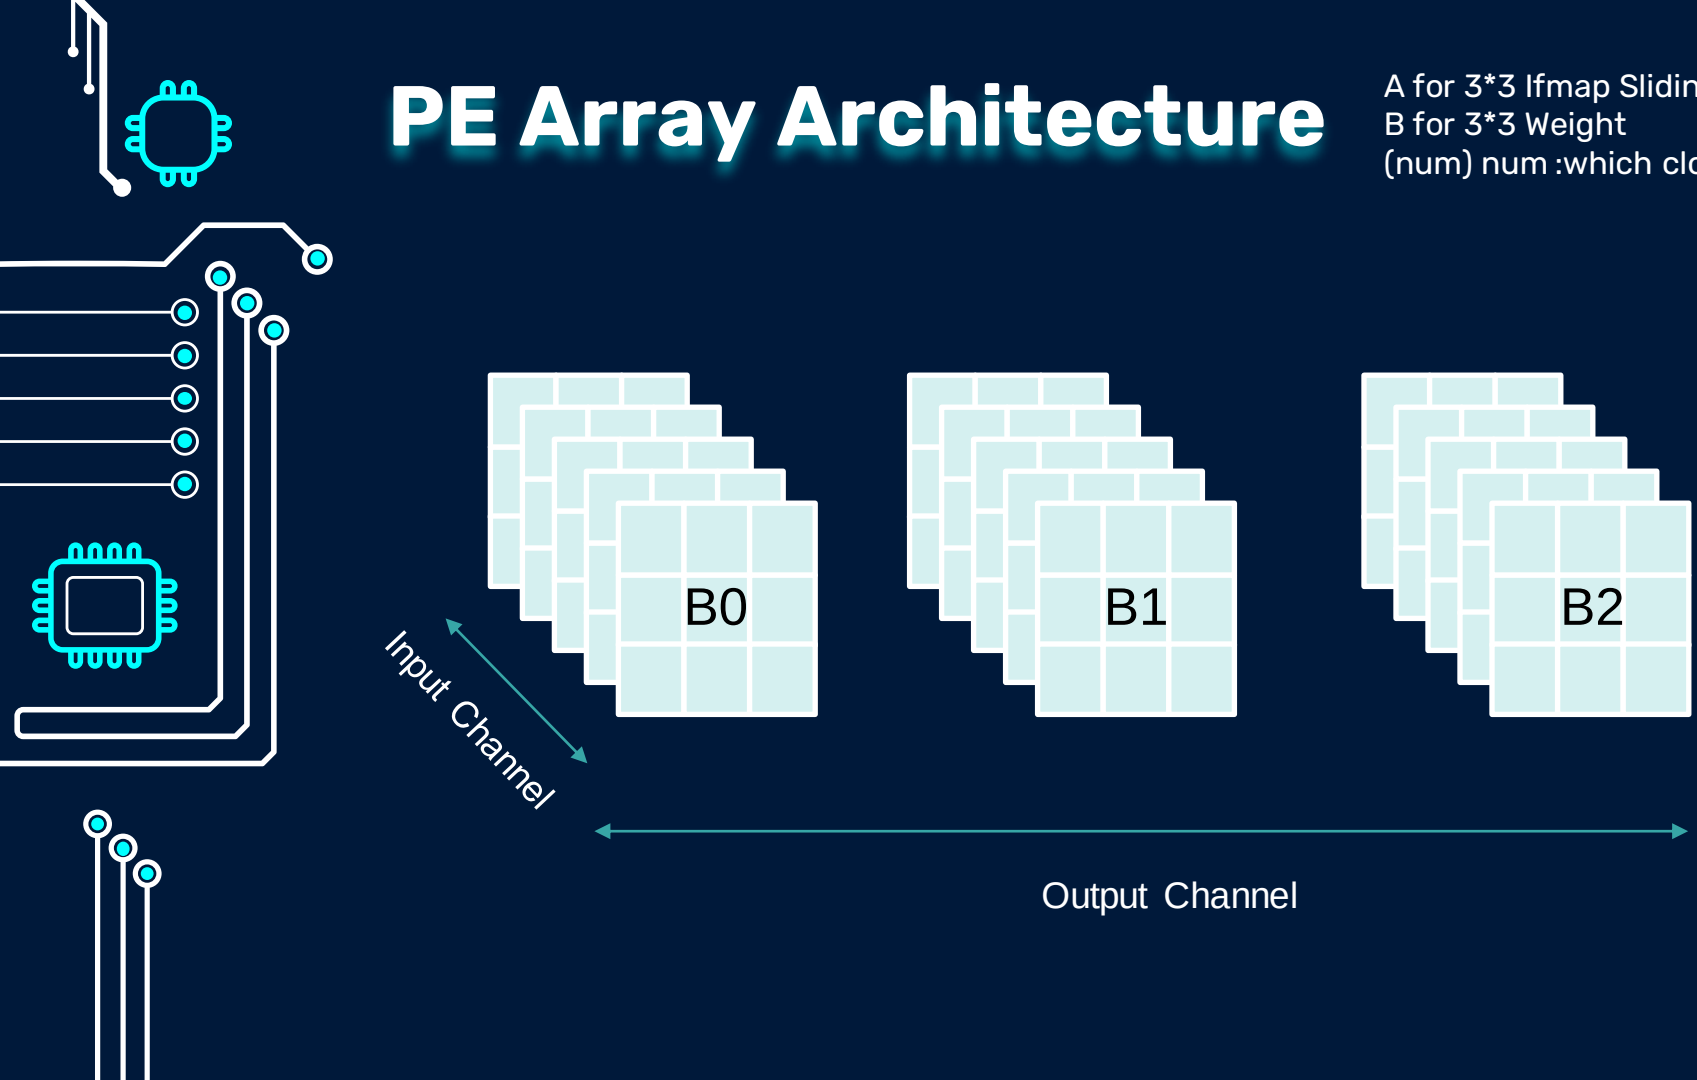

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

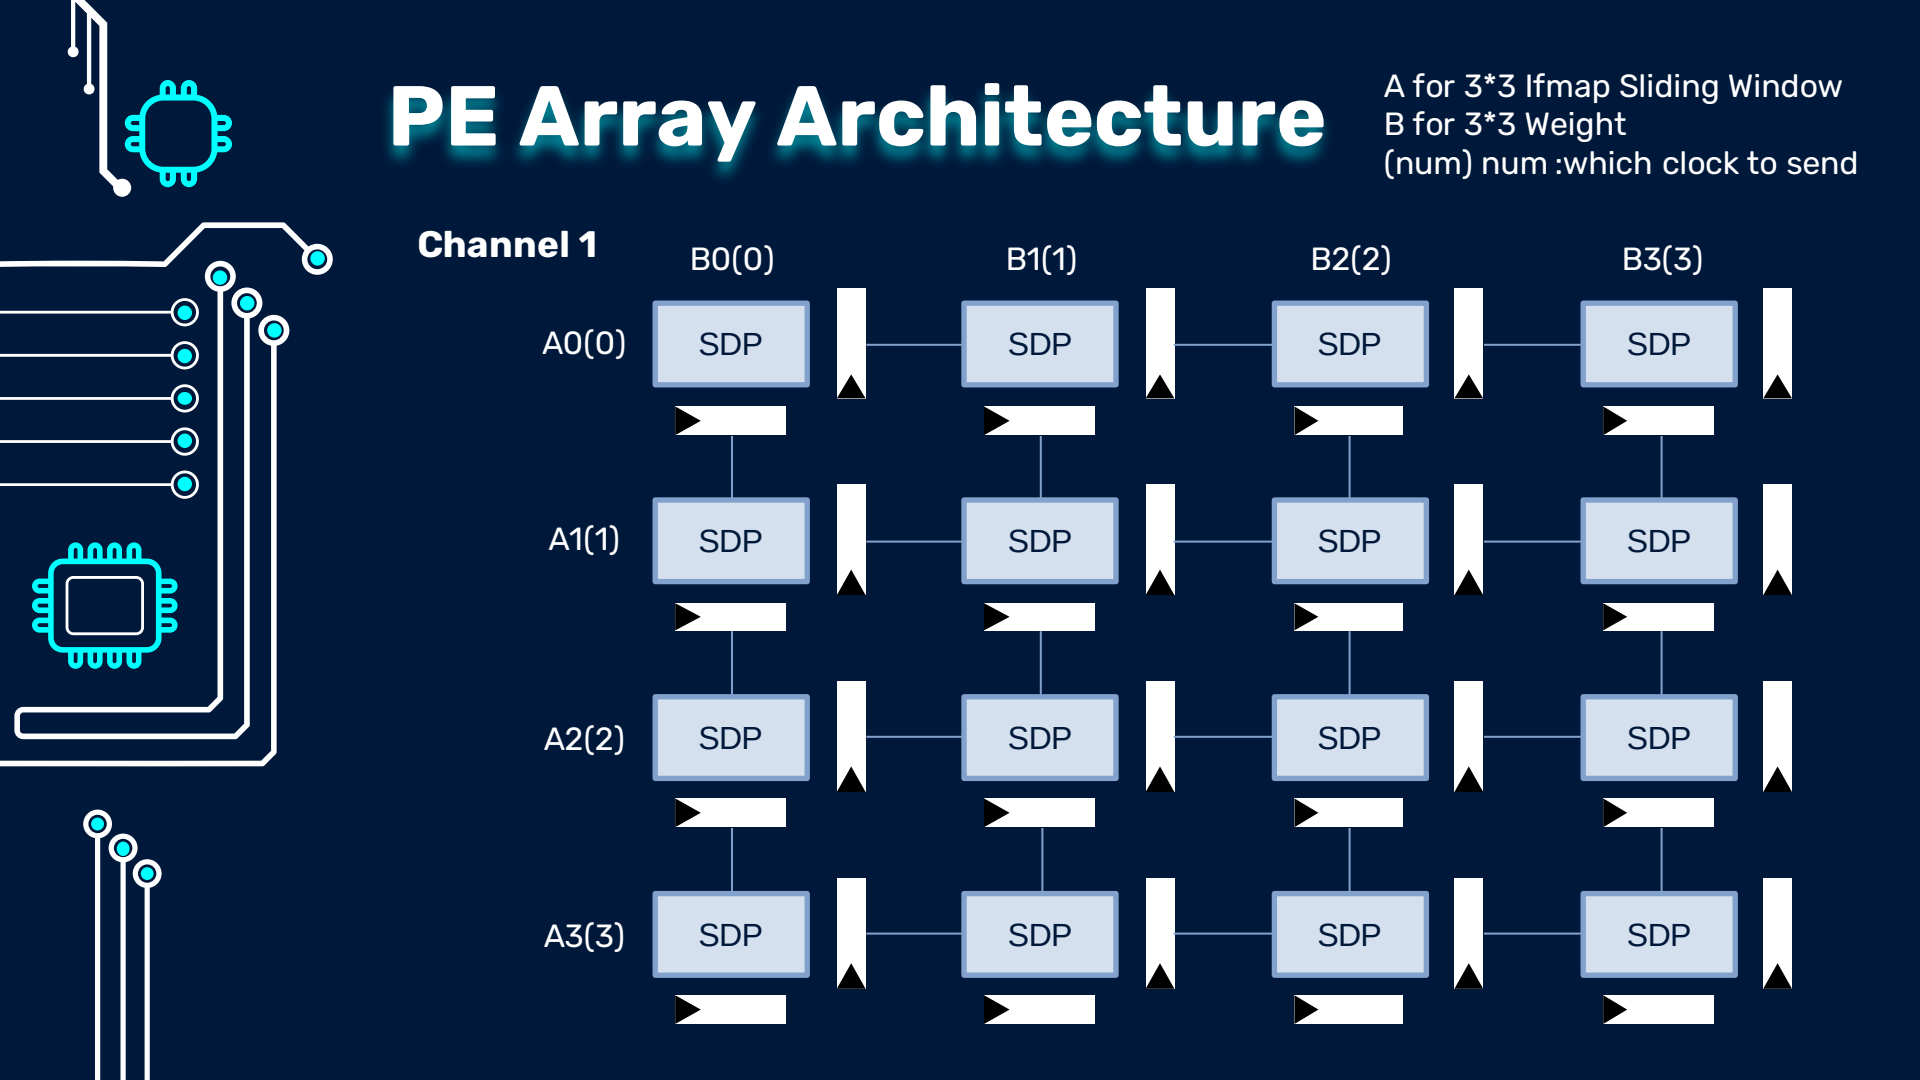

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

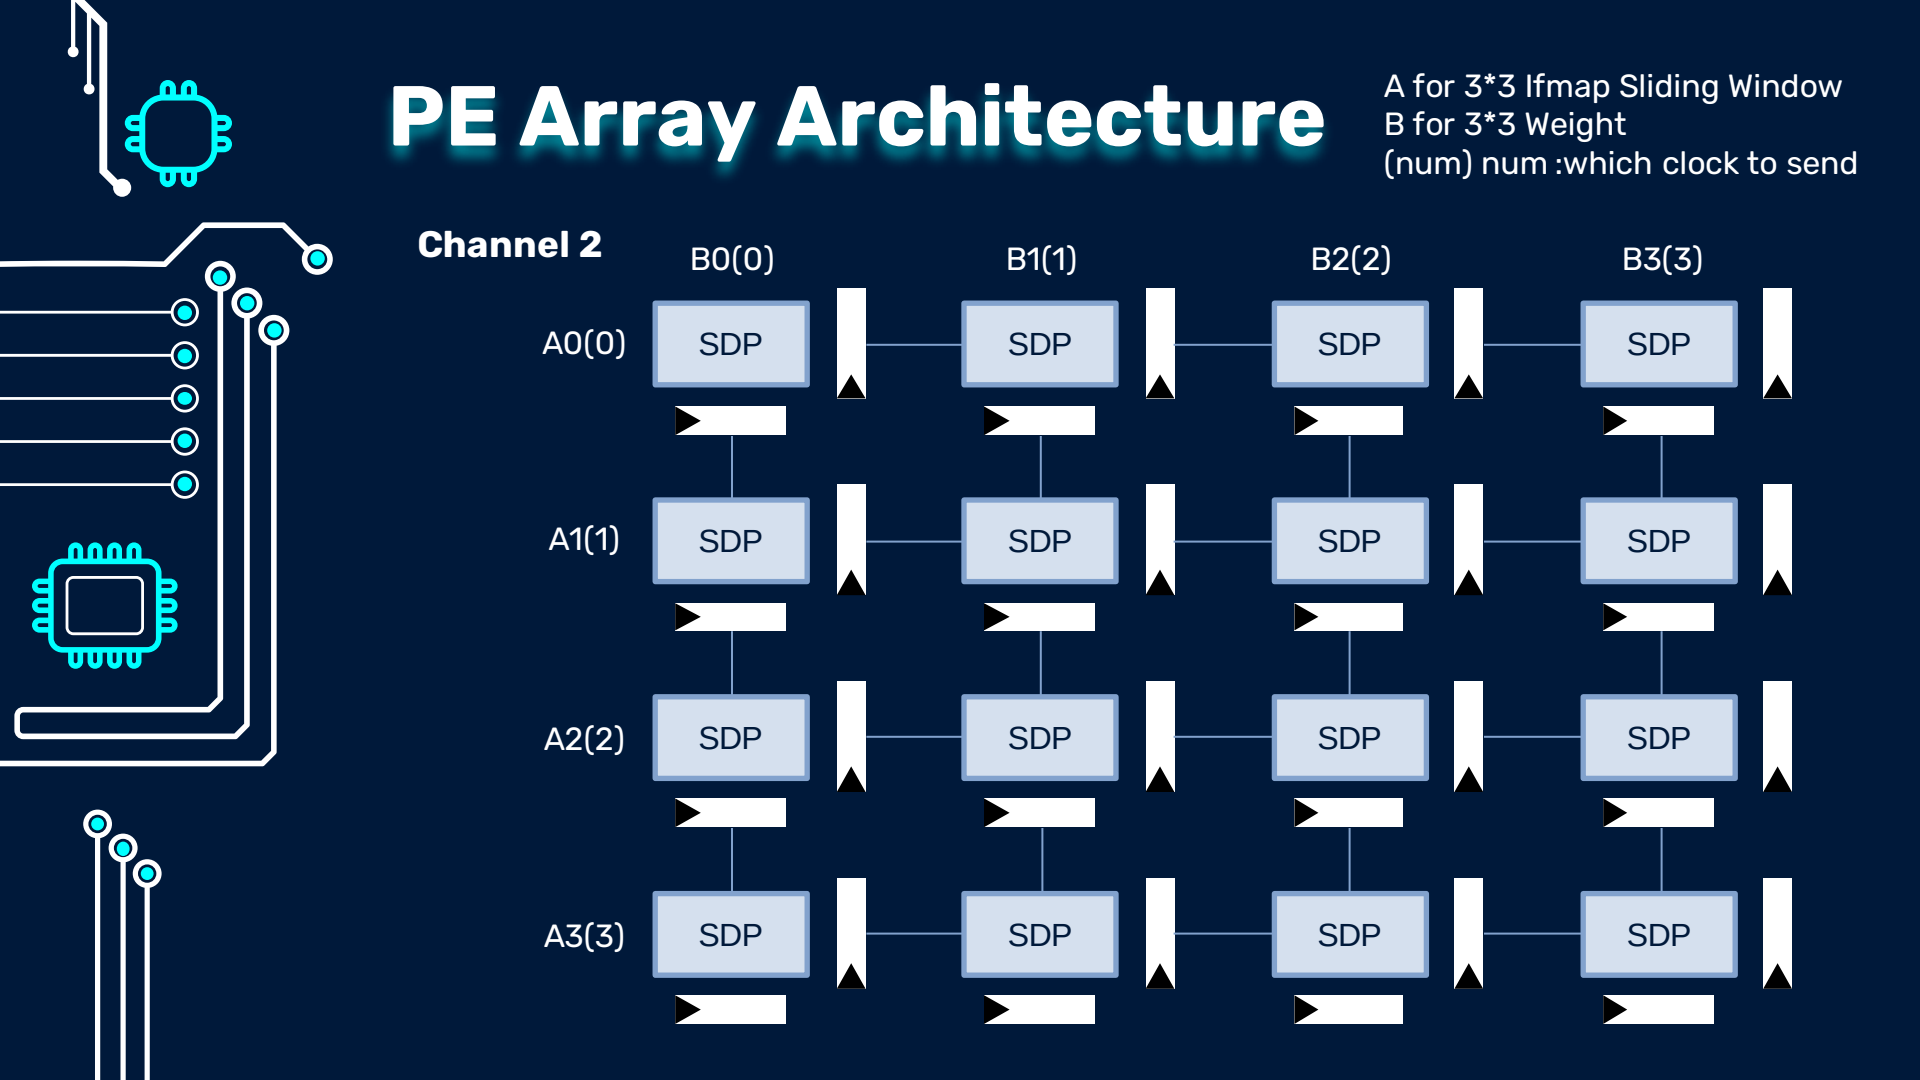

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

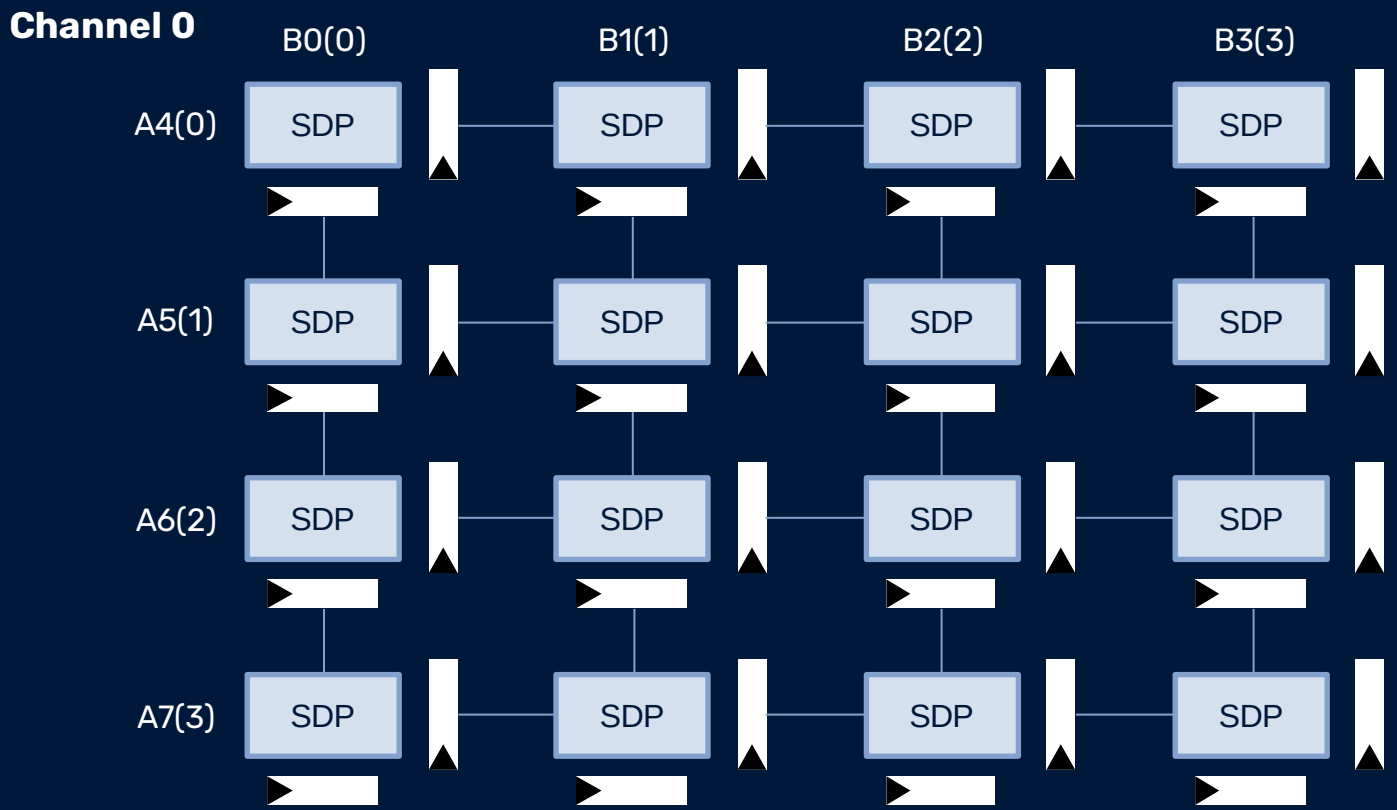

PE Array Architecture

A for 3*3 Ifmap Sliding Window
B for 3*3 Weight
(num) num :which clock to send

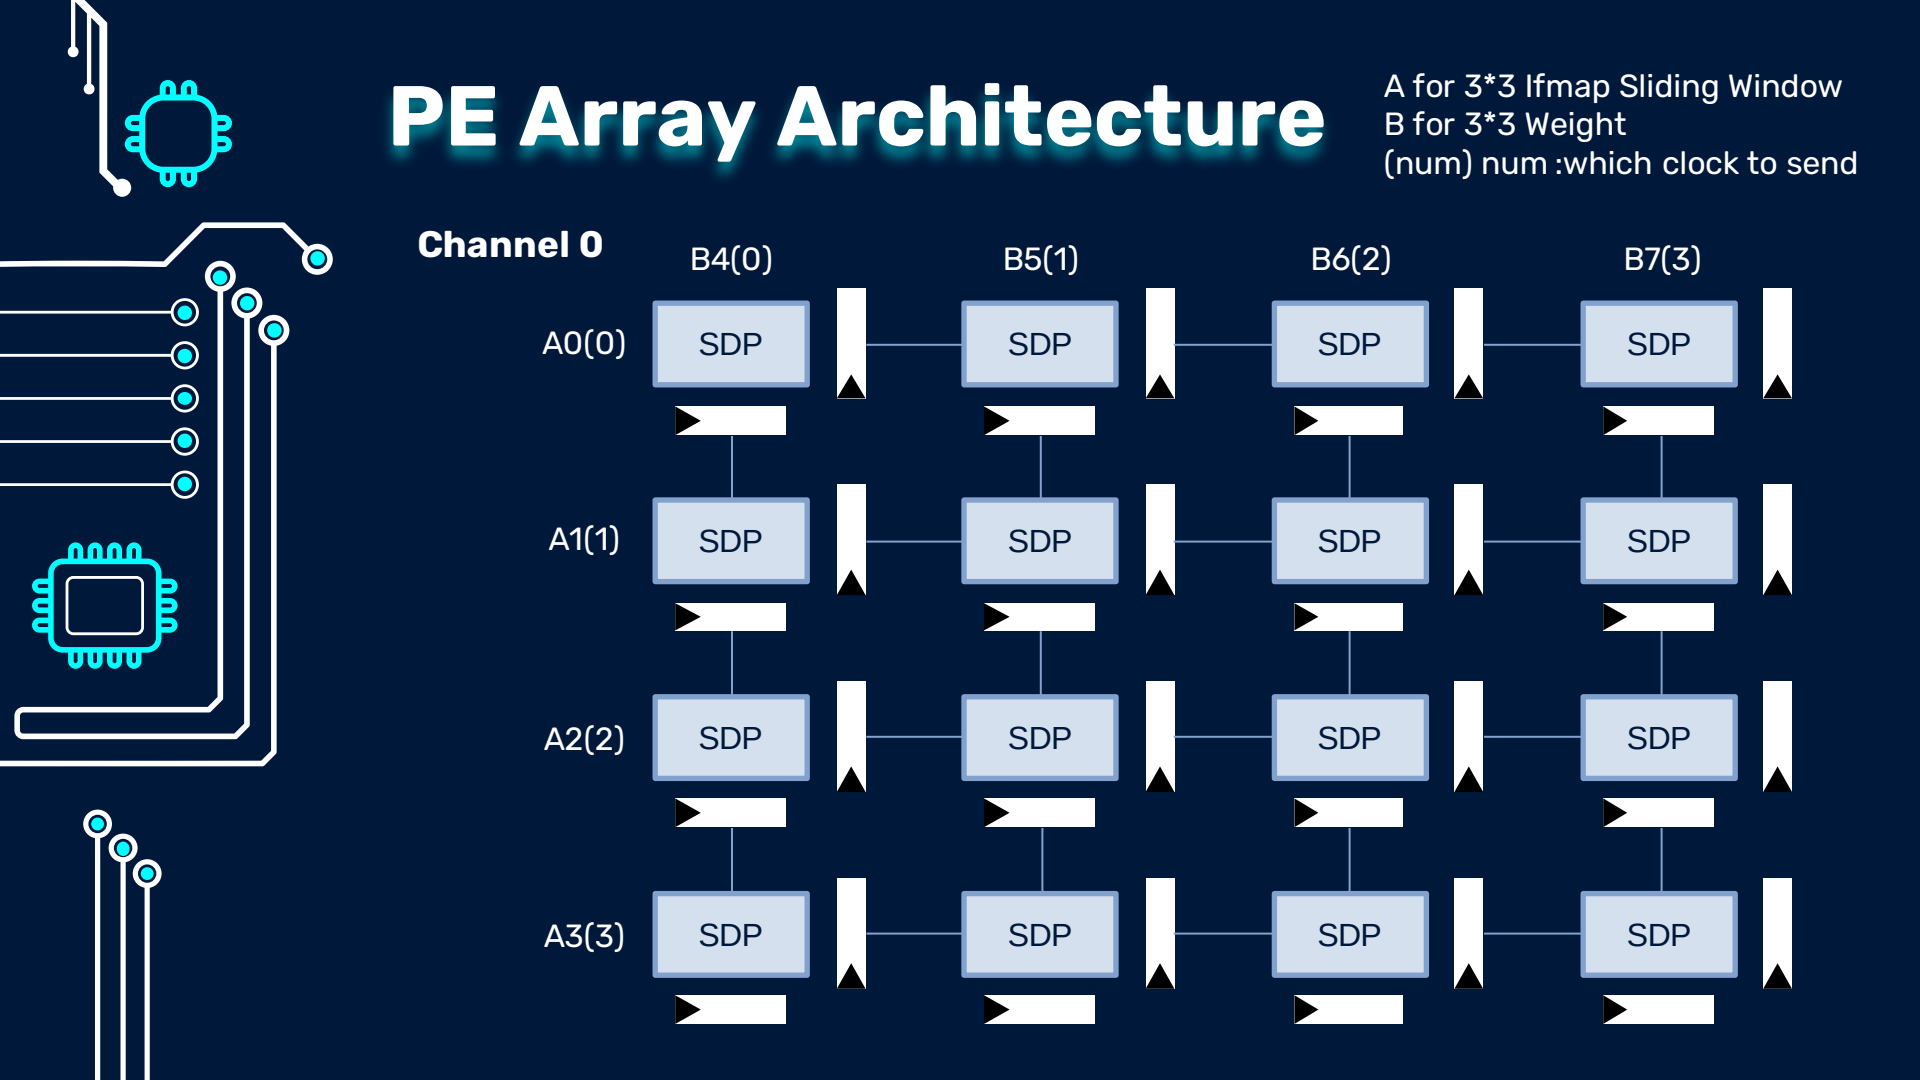

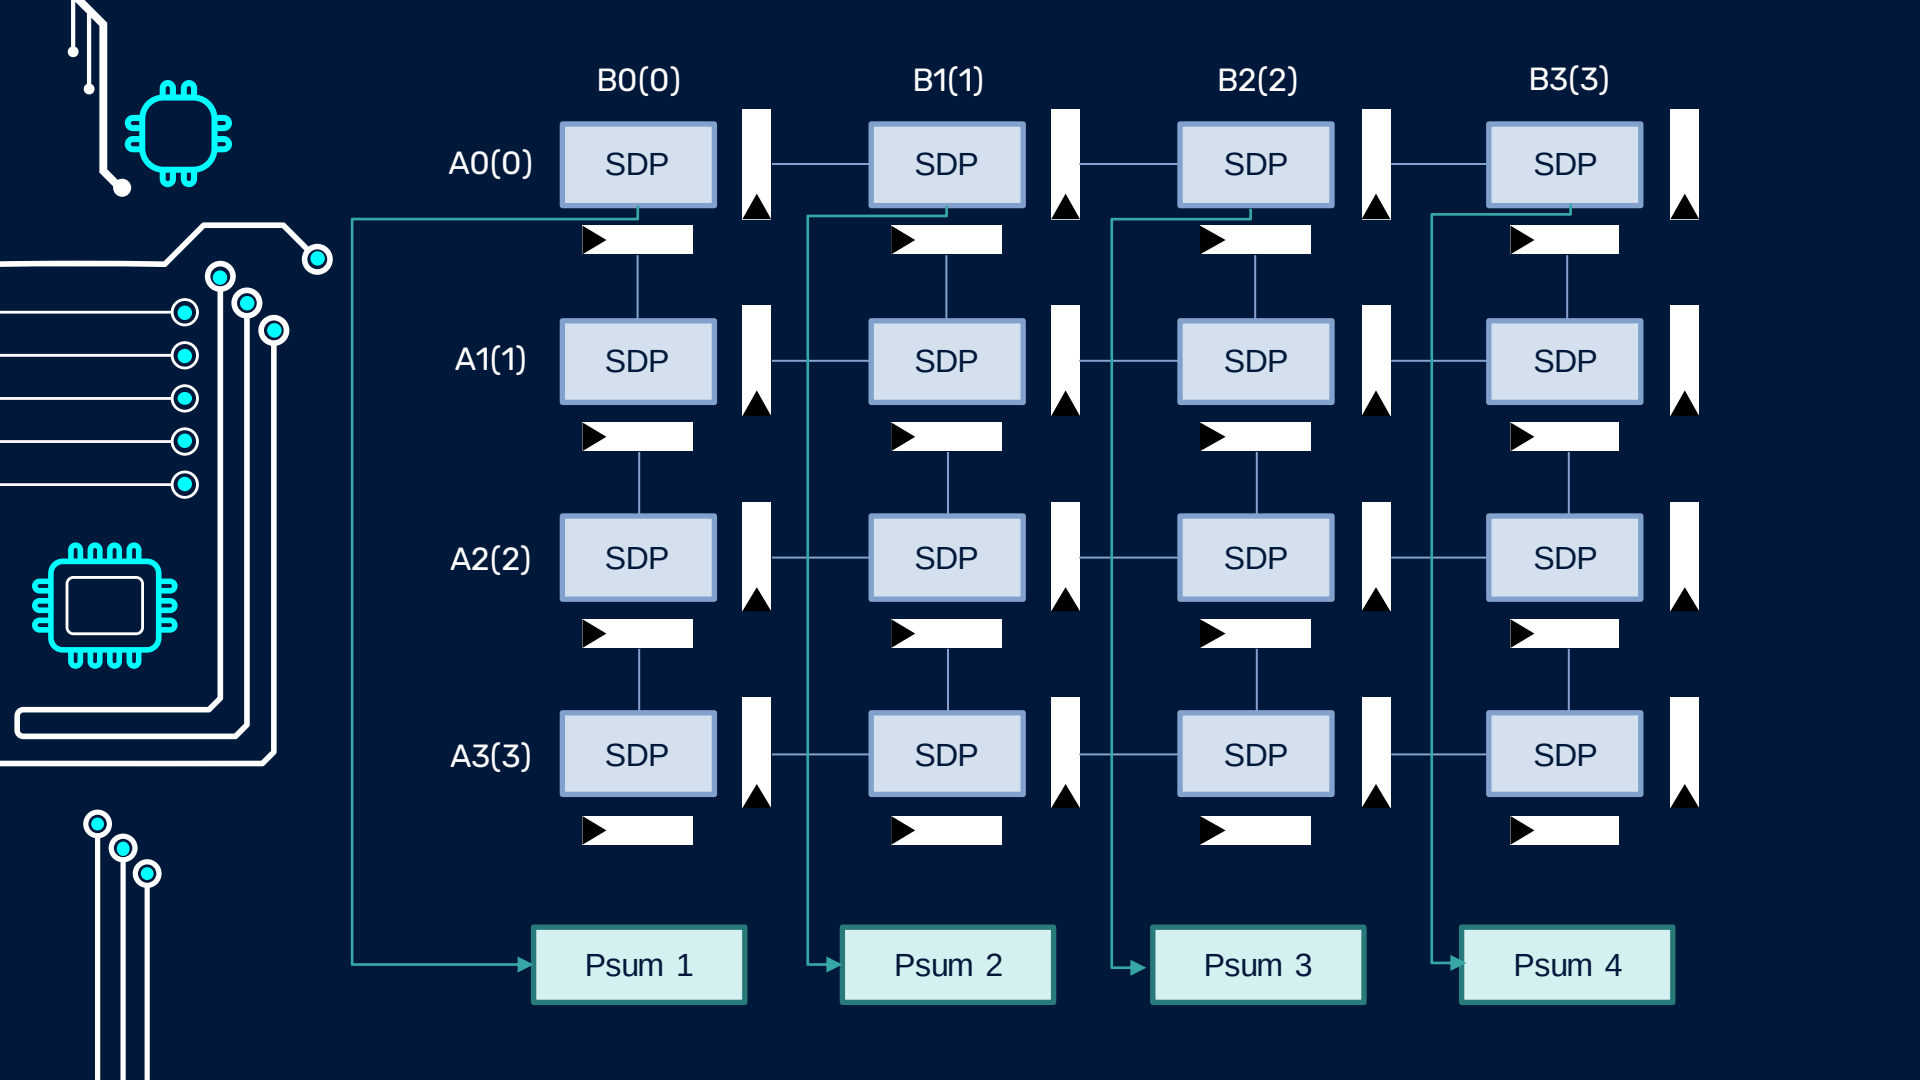
B0(0)
B1(1)
B2(2)
B3(3)
A0(0)
A1(1)
A2(2)
A3(3)
SDP
Psum 1
Psum 2
Psum 3
Psum 4

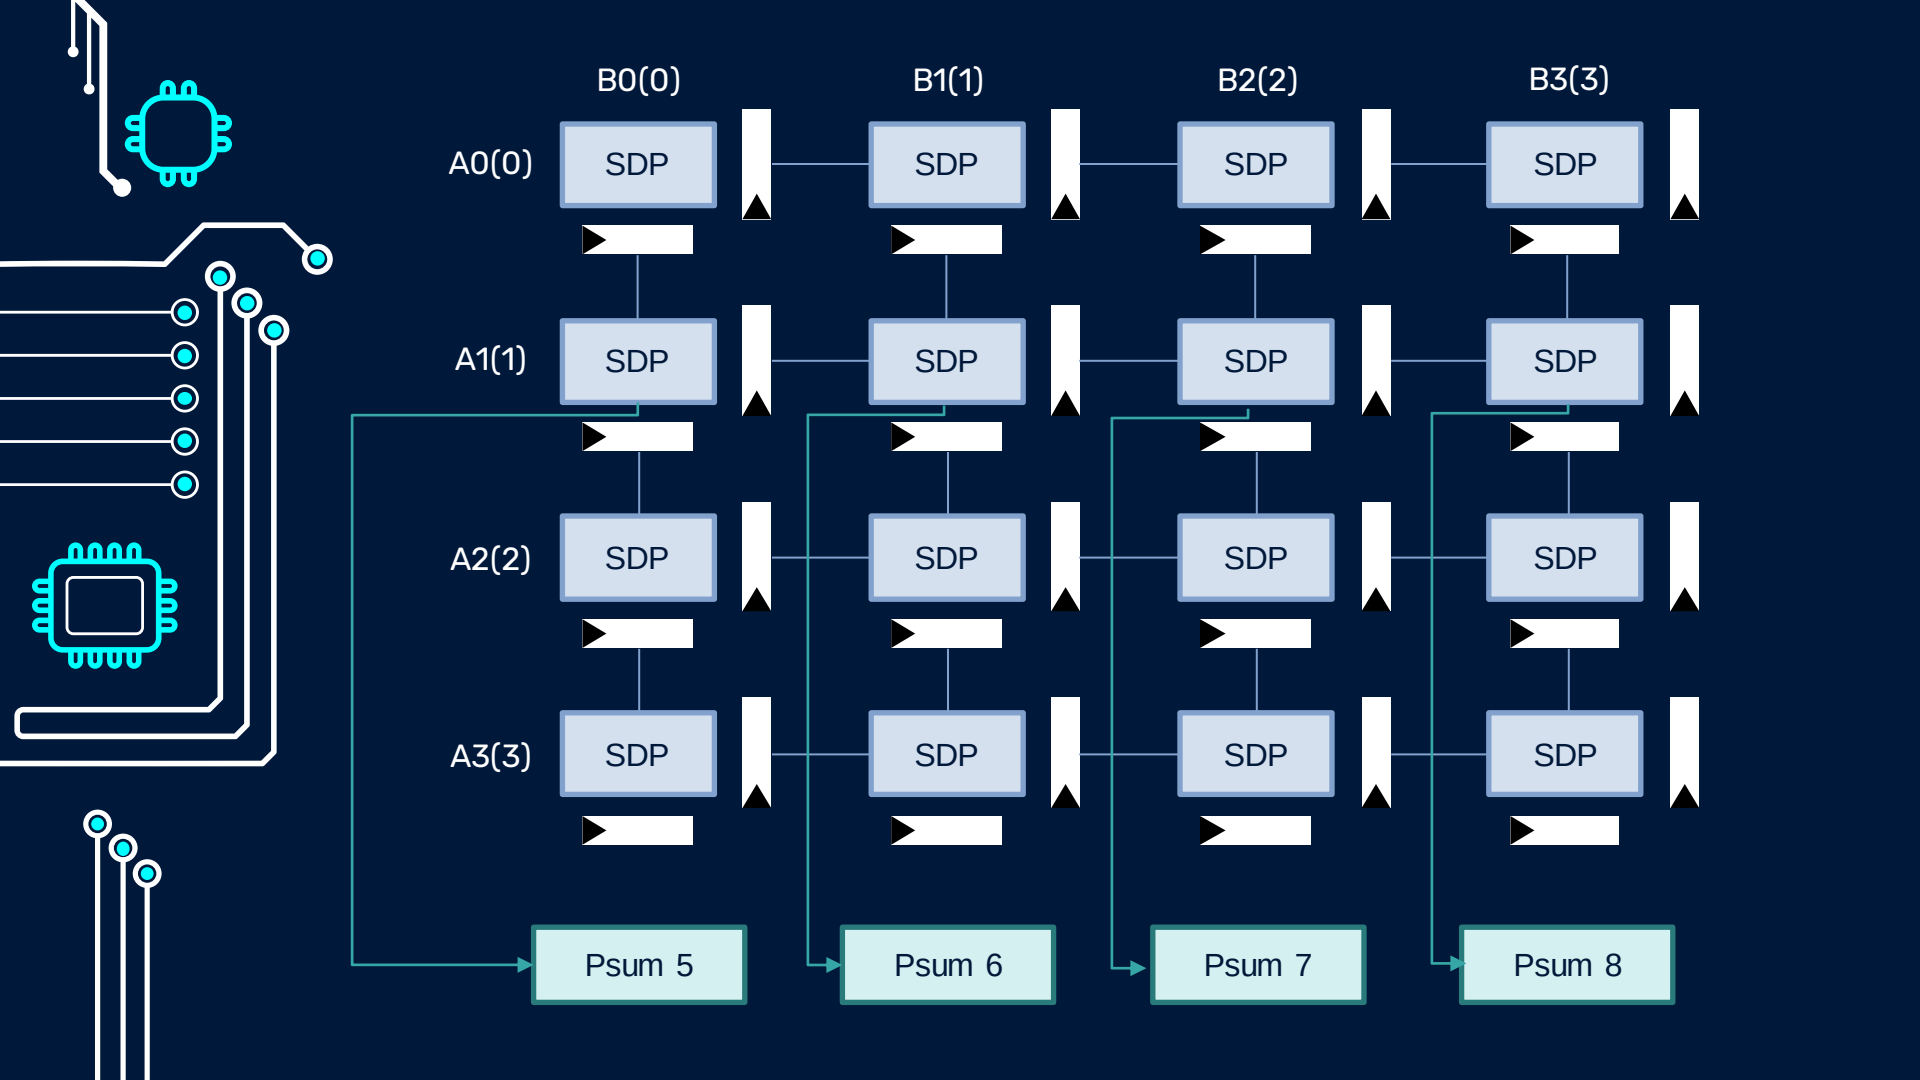
B0(0)
B1(1)
B2(2)
B3(3)
A0(0)
A1(1)
A2(2)
A3(3)
SDP
Psum 5
Psum 6
Psum 7
Psum 8

03

Our Accelerator Architecture and SPEC

Architecture

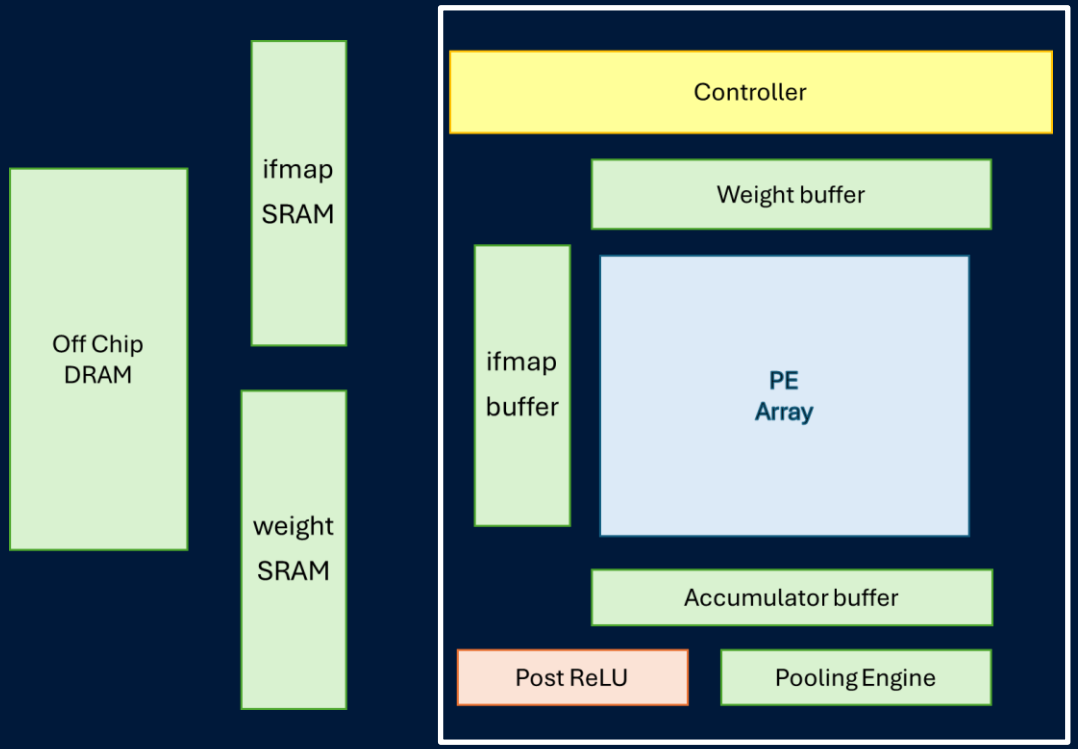

SPEC

SRAM

Type	Size	Why we use this size?
Ifmap SRAM	16KB	512(input channel) * 3(filter row) * 10(6+4) * 1Byte = 15KB
Weight SRAM	16KB	512(input channel) * 8(output channel) * 3(NNZ) * 10(int8+2 bit address) = 15KB

SPEC

On-Chip Buffer

Type	Size	Why we use this size?
Ifmap Buffer	576B	32(input channel) * 3(filter row) * 6(ifmap column) * 1Byte = 576B
Weight Buffer	480B	32(input channel) * 3(filter NNZ) * 4(PE a row) * 10bit (int8 + 2 bit address) = 480B
Accumulator Buffer	64B	32 Bits * 16 (PE) = 64B

SPEC

Bandwidth

Type	Size	Why we use this size?
PE to Ifmap Buffer	28B	7(element)*4(PE a row)*1Byte = 28B
PE to Weight Buffer	15B	3(NNZ in a filter) * 4(PE a col) * 10 bits = 15B
PE to Accumulator Buffer	64B	4Byte * 16 = 64B

04

Analysis and Question

Advantages

- Increased Efficiency
- Enhanced Area and Power Efficiency
- Efficient Data Flow

Roofline Model

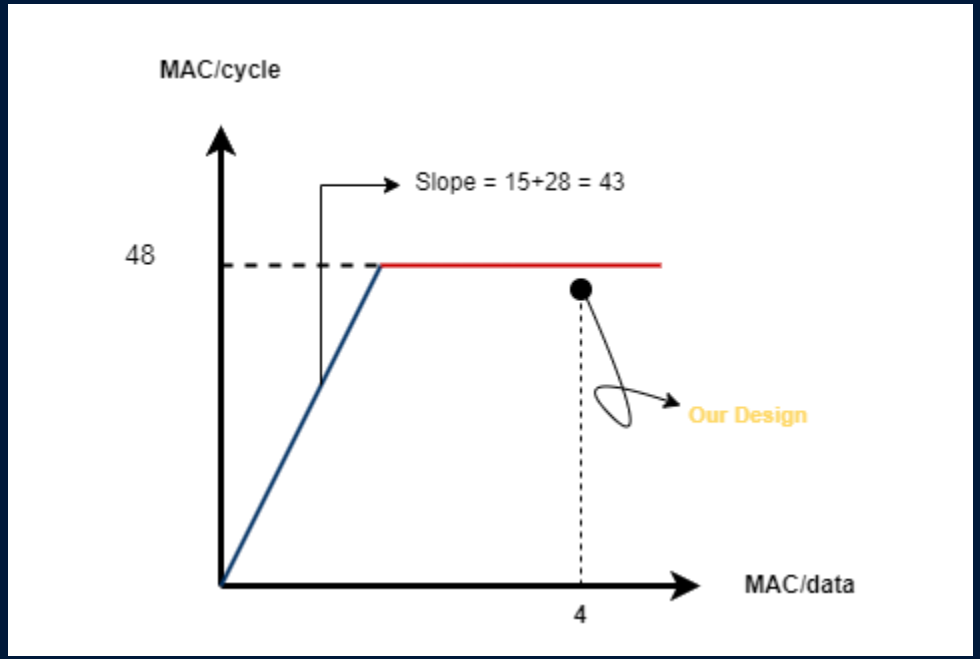

Problem Disscusion

Q1: Is the Accumulator on-chip bandwidth (64B) too large?

Q2: Is the number of times data is fetched from DRAM too many?

(Up to 64 times per layer)

RESOURCES

Course

- AI on Chip NCKU by Professor Chia-Chi Tsai

Paper Reference

- **Liu, Z. G., Whatmough, P. N., & Mattina, M. (2020). Systolic tensor array: An efficient structured-sparse GEMM accelerator for mobile CNN inference. IEEE Computer Architecture Letters, 19(1), 34-37.**
- **Niu, W., Ma, X., Lin, S., Wang, S., Qian, X., Lin, X., ... & Ren, B. (2020, March). Patdnn: Achieving real-time dnn execution on mobile devices with pattern-based weight pruning. In Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems (pp. 907-922).**

THANKS!