

AOC Design Proposal

Group 2 軟硬兼施

施宇庭 NN6124030 陳均嫻 NP8124014

葉宣佑 E24096695 王力恩 E24091108 施尚甫 E64096148

Targeting Tasks

- Image Classification on Embedded System

- GoogLeNet

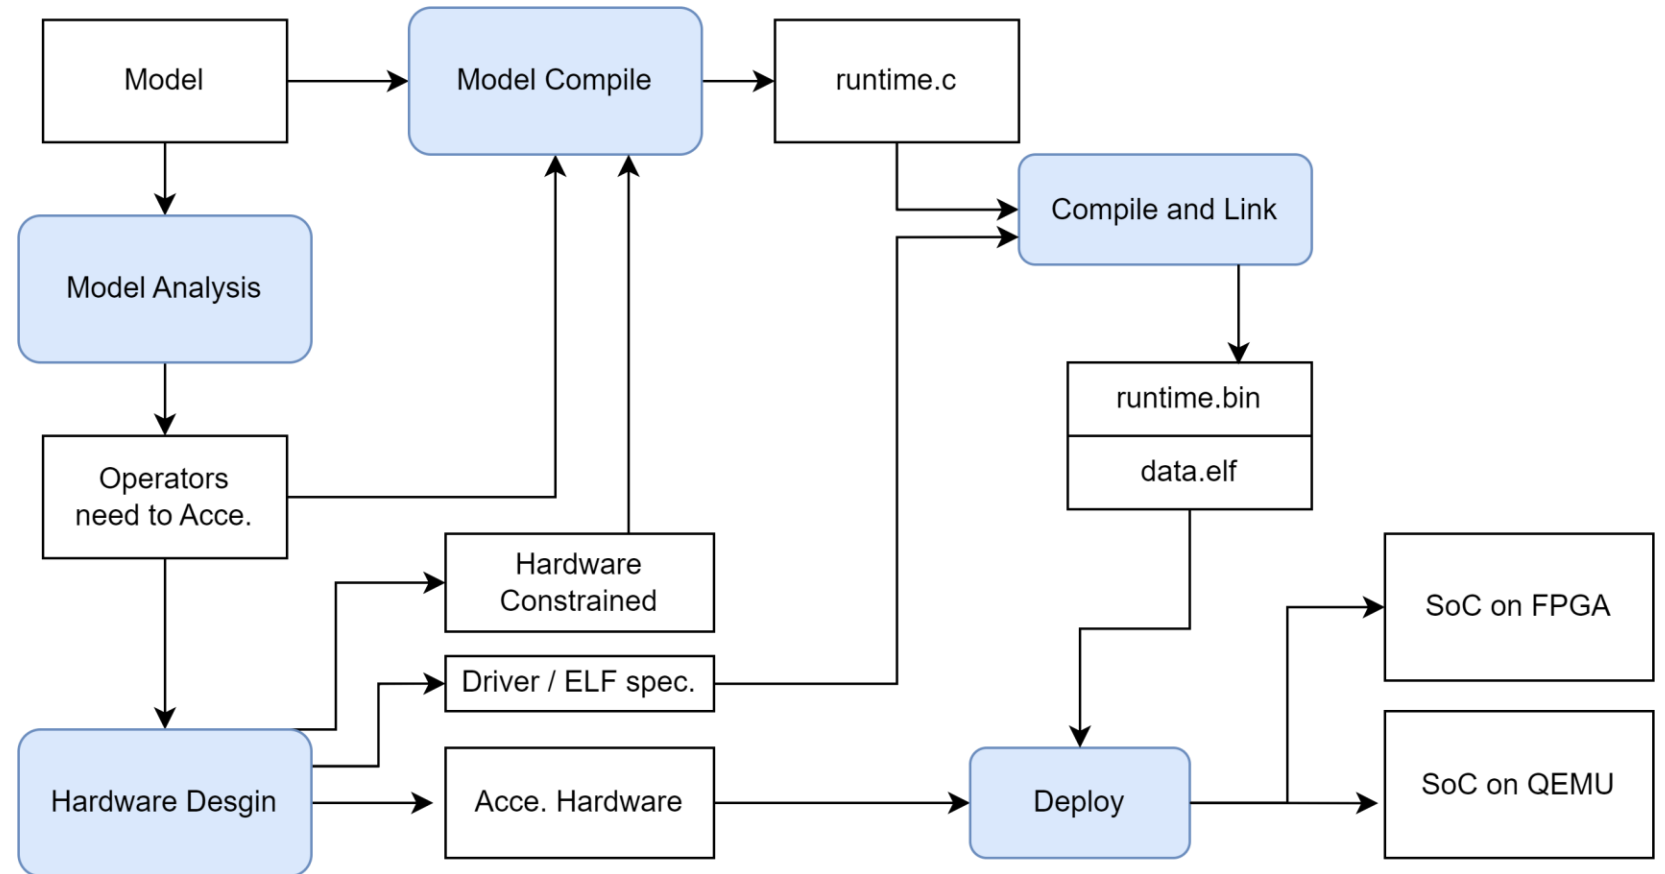
- Operator - Conv2D, MaxPool
 - Parameters - 6.6M
 - model size - 27MB / 7MB

- SqueezeNet

- Operator - Conv2D, MaxPool2D, BatchNorm, Concat, GlobalAvgPool2D, Softmax
 - Parameters - 1.2M
 - model size - 5MB / 2MB

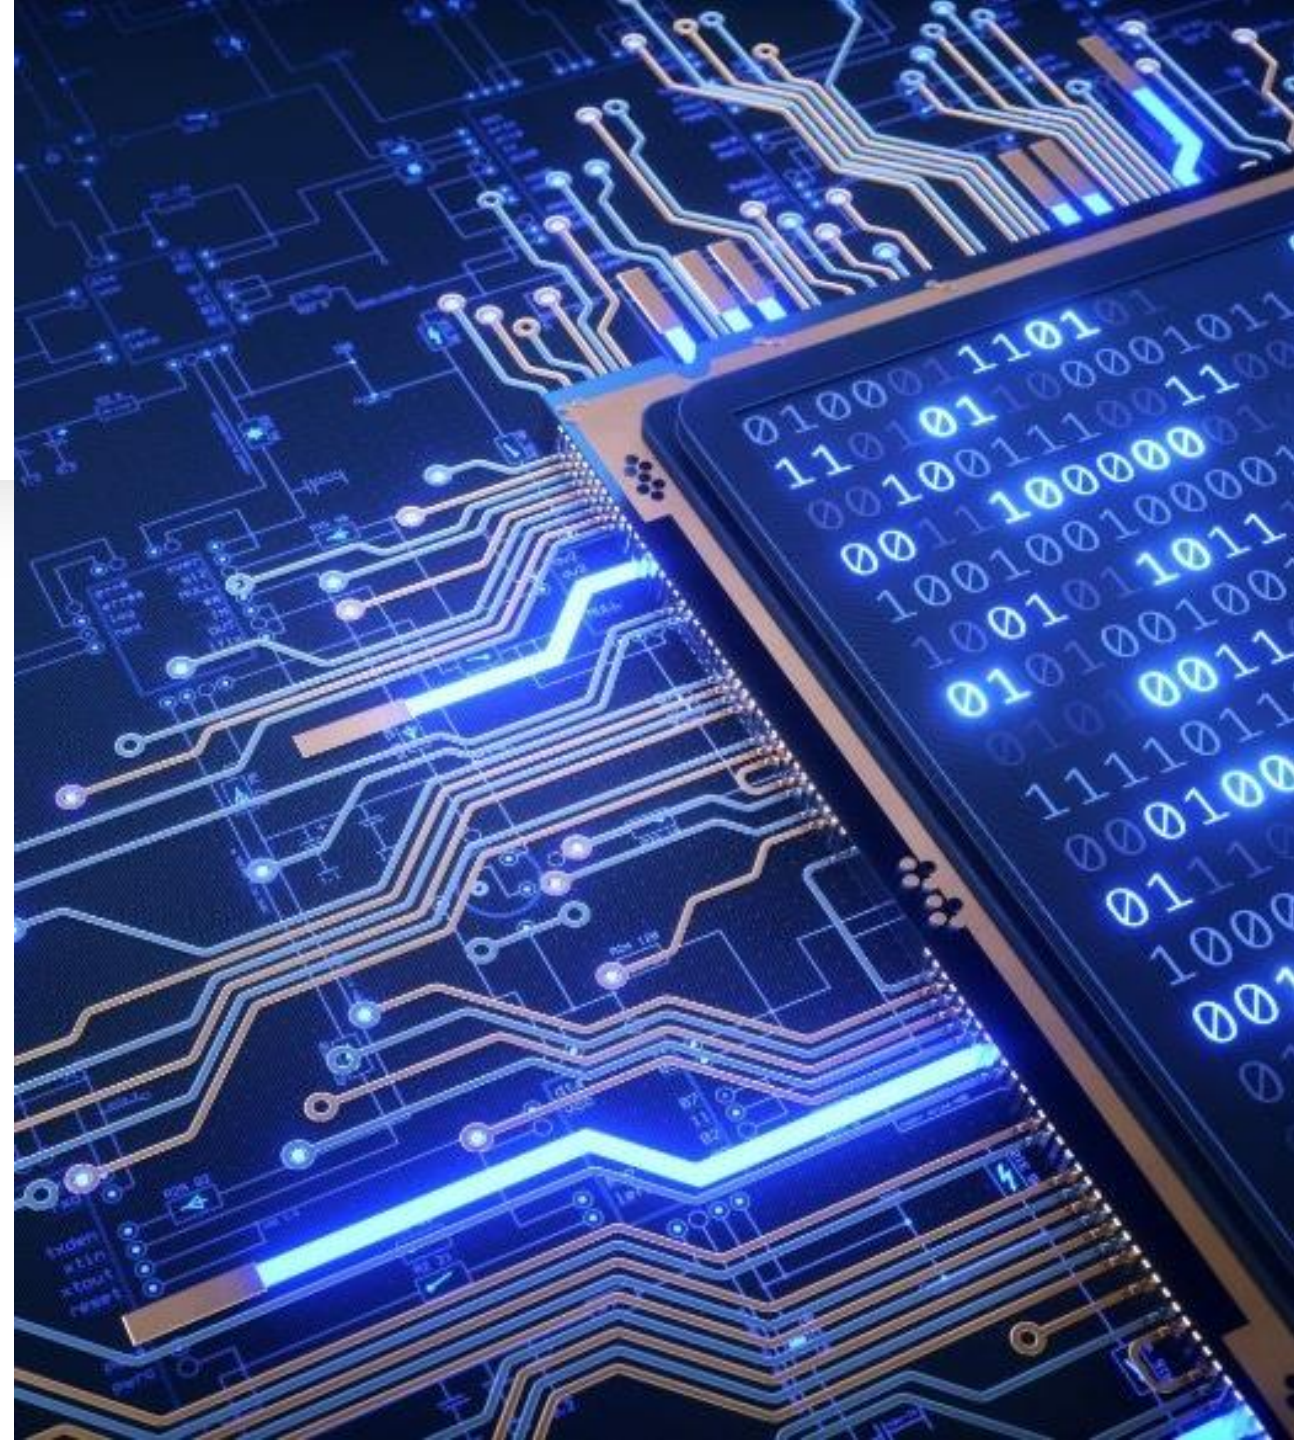
- Machine Comprehension
 - BERT-Squad 均燧
 - BERT-Squad (416 MB)
 - BERT-Squad-int8 (119 MB)
 - BiDAF
 - BiDAF (41.5 MB)
 - BiDAF-int8 (12 MB)
- Image Classification
 - MobileNet 均燧 Yoyo Yeh
 - MobileNet v2-1.0-fp32 (13.3 MB)
 - MobileNet v2-1.0-int8 (3.5 MB)
 - ResNet
 - ResNet50-fp32 (97.0 MB)
 - ResNet50-int8 (24.6 MB)
 - SqueezeNet 均燧 yuting
 - SqueezeNet 1.0 (5 MB)
 - SqueezeNet 1.0-int8 (2 MB)
 - VGG
 - VGG-16-fp32 (527.0 MB)
 - VGG-16-int8 (132.0 MB)
 - AlexNet
 - AlexNet (233 MB)
 - AlexNet-int8 (58 MB)
 - GoogleNet 均燧 yuting
 - GoogleNet (27 MB)
 - GoogleNet-int8 (7 MB)
 - CaffeNet
 - CaffeNet (233 MB)
 - CaffeNet-int8 (58 MB)
 - DenseNet-121 均燧 yuting
 - DenseNet-121-12 (32 MB)
 - DenseNet-121-12-int8 (9 MB)
 - Inception
 - Inception-1 (27 MB)
 - Inception-1-int8 (10 MB)
 - ☑ ShuffleNet 均燧 Yoyo Yeh
 - ShuffleNet-v2-fp32 (8.79MB)
 - ShuffleNet-v2-int8 (2.28MB)
 - ZFNet-512 均燧 均燧
 - ZFNet-512 (333 MB)
 - ZFNet-512-int8 (83 MB)
 - EfficientNet-Lite4 均燧 均燧
 - EfficientNet-Lite4 (51.9 MB)
 - EfficientNet-Lite4-int8 (13.0 MB)
 - YOLOv7 均燧 力悉

Design Flow



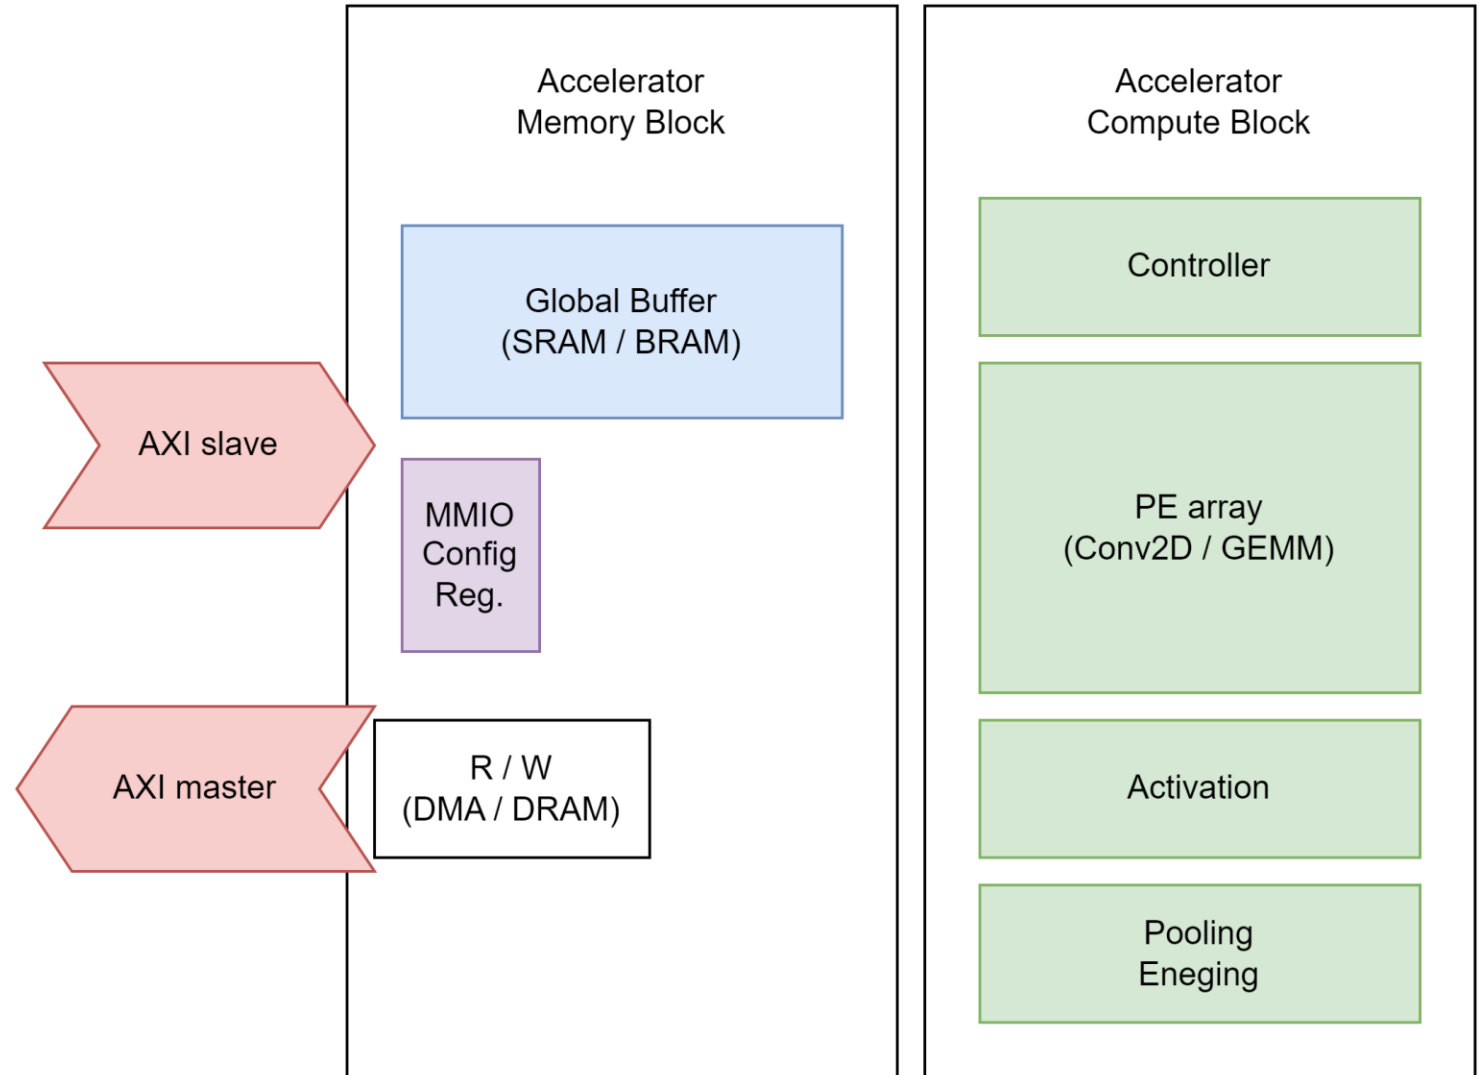
Hardware Design

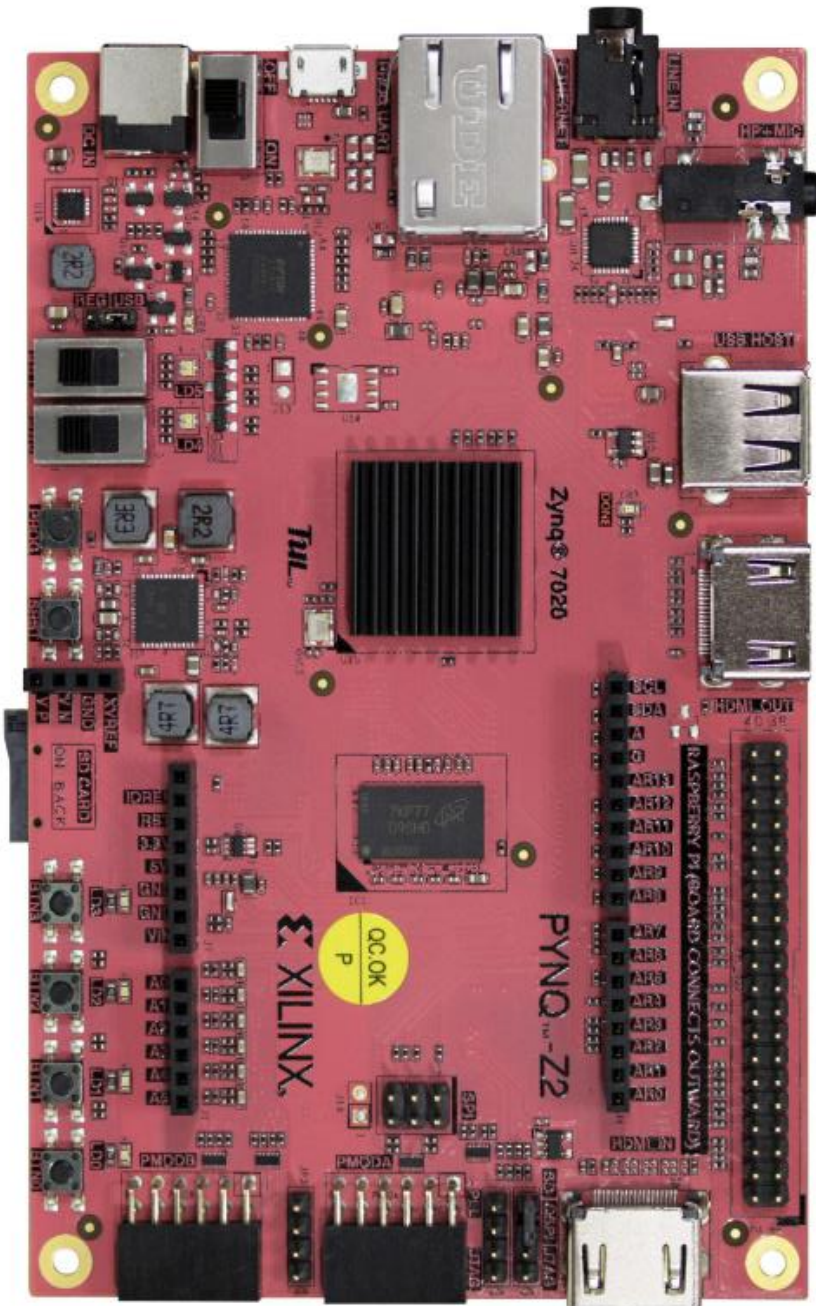
- Architecture
- Verification



Architecture

- Design based on
 - NVDLA
 - Eyeriss
- Support Operation
 - Conv2D
 - GEMM
 - MaxPool2D
 - AvgPool2D
 - ReLU
 - Softmax



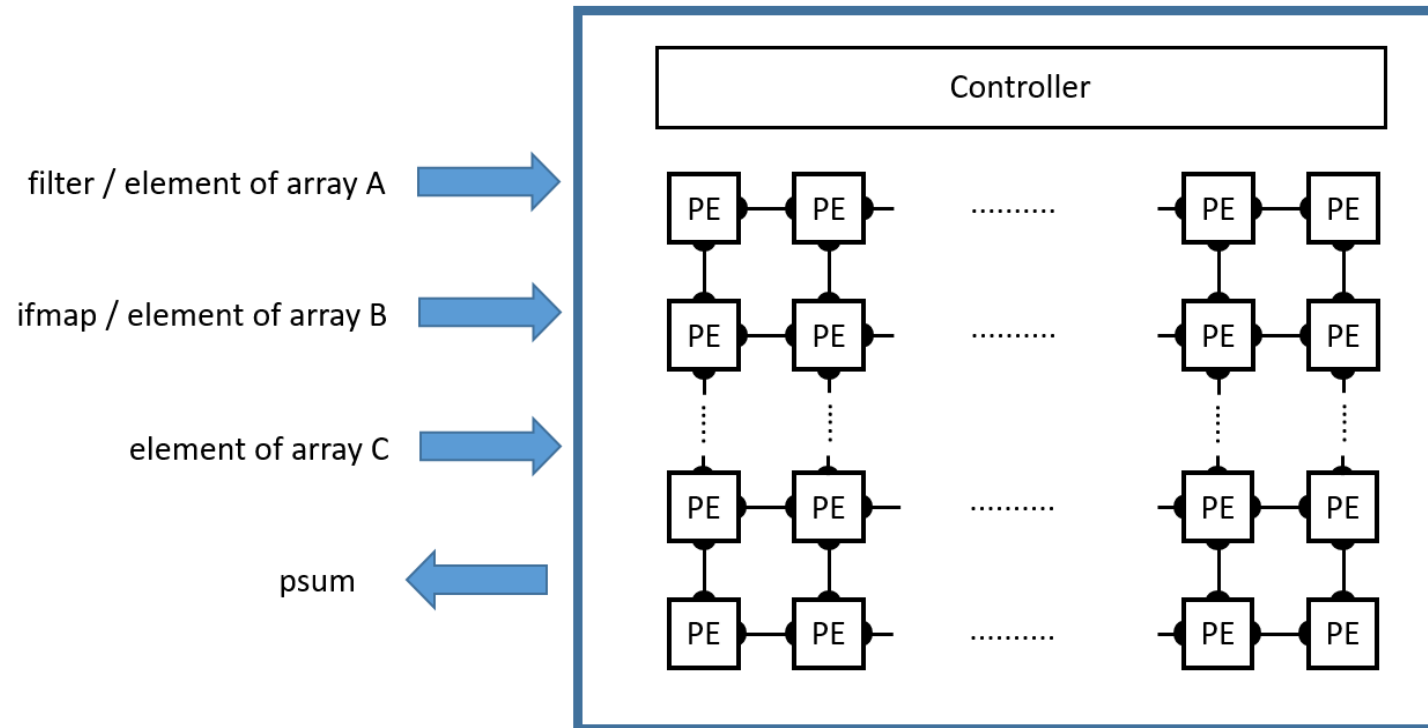


Verification Platform

- Plan A – Porting on FPGA
 - PYNQ-Z2
 - RAM512MB
 - Pros: cycle accurate
 - Cons: synthesis technical problem (AXI / BRAM)
- Plan B – QEMU with C simulation
 - QEMU with custom virtual device
 - Pros: fast deploy
 - Cons: only behavior simulations

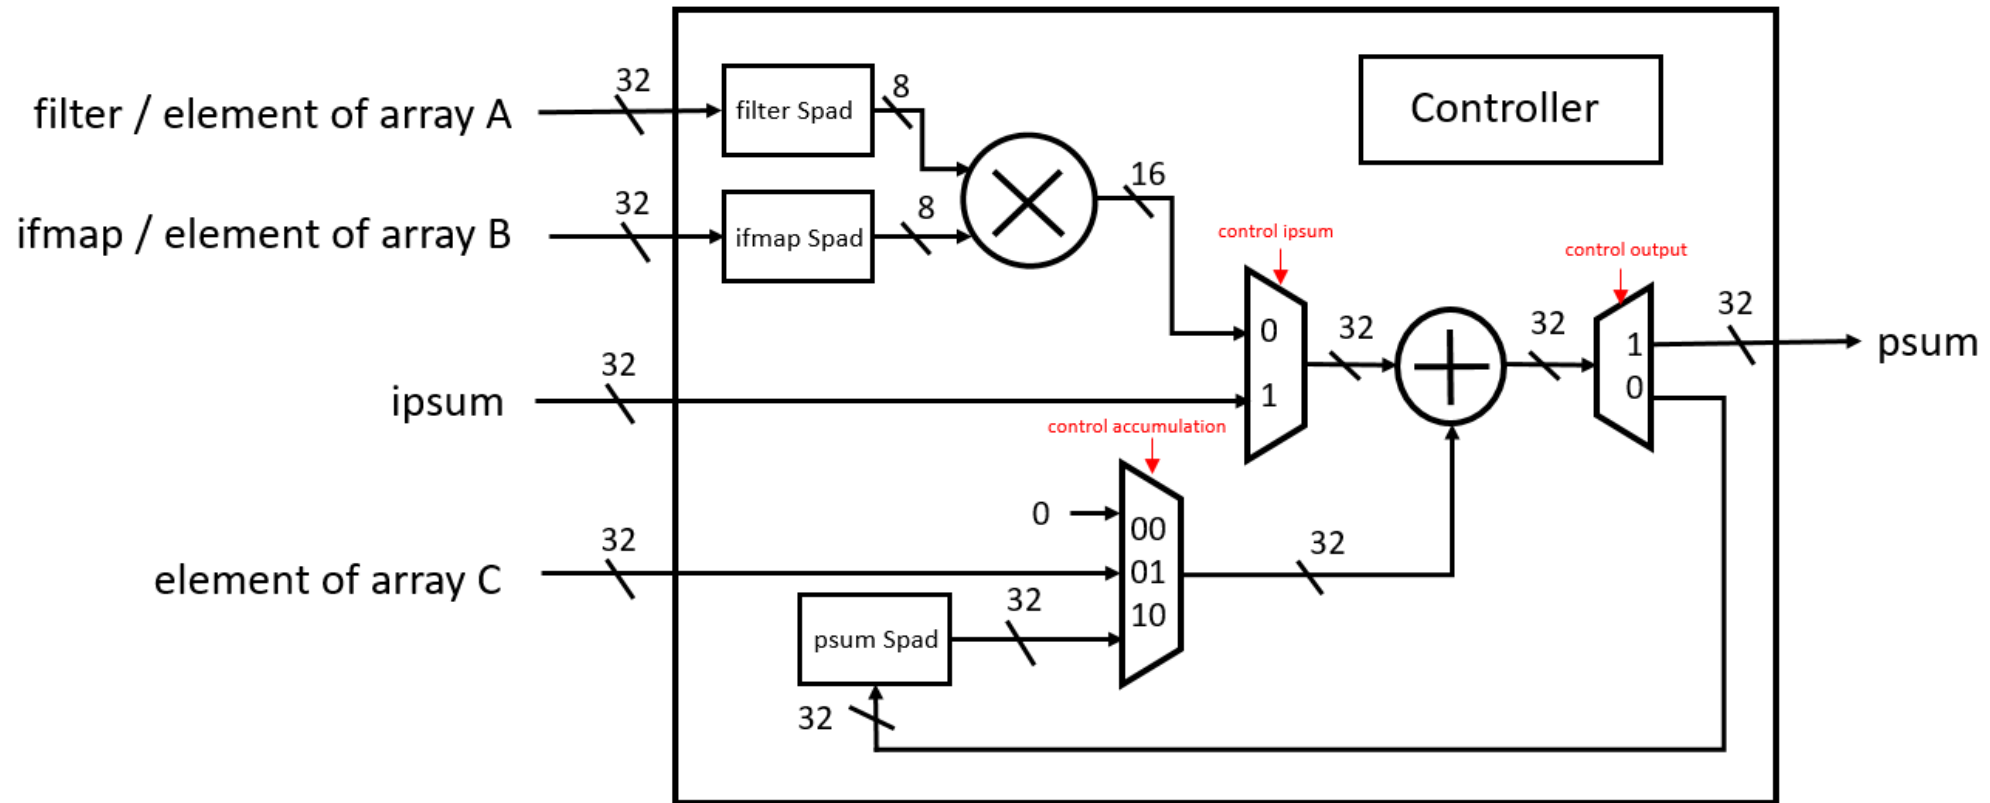
PE array

- Conv2D (1x1, 3x3, 5x5, 7x7), GEMM
- Size of PE array will be determined by the supported Conv2D size



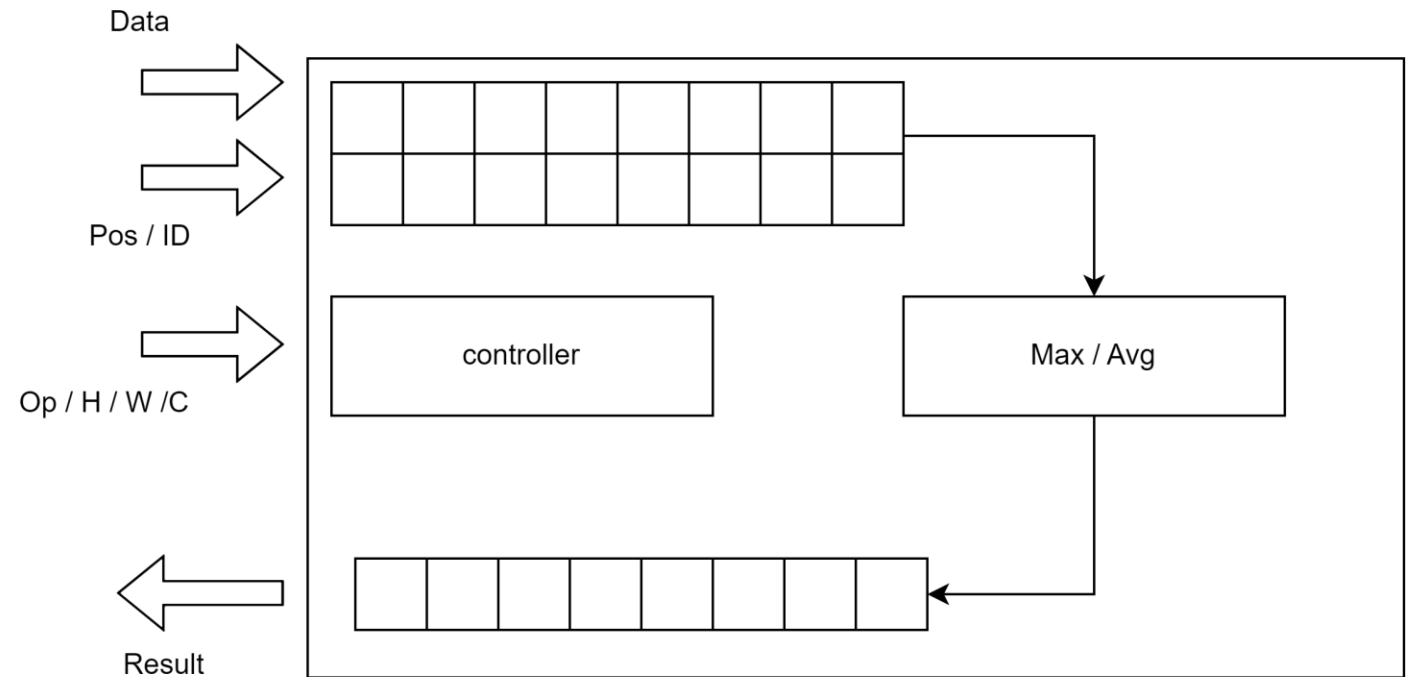
PE Architecture

- Scratchpad size will be determined by the result of analytical model



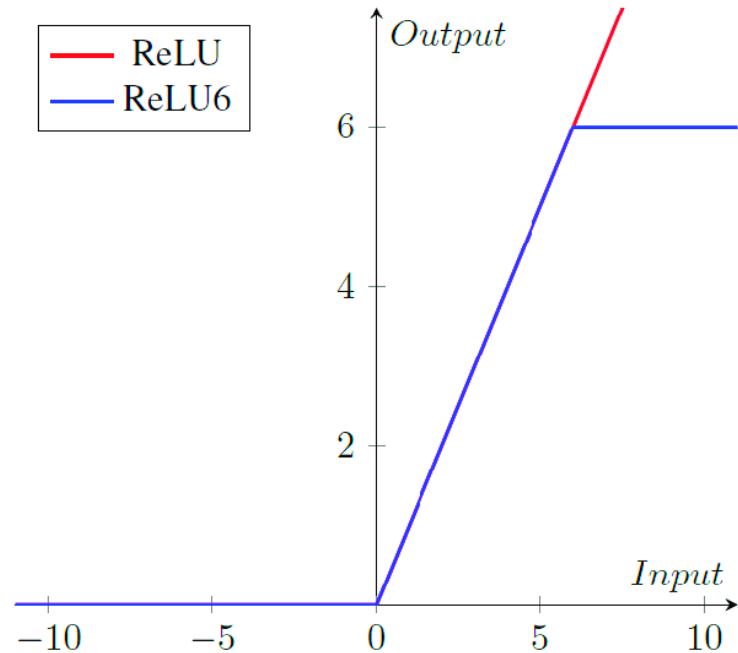
Pooling Engine

- MaxPooling
 - (2x2, stride=(2,2))
 - (3x3, stride=(1,1))
 - (3x3, stride=(2,2))
- AvgPooling
 - GlobalAveragePooling2D



Activation Engine

- ReLU
- ReLU6



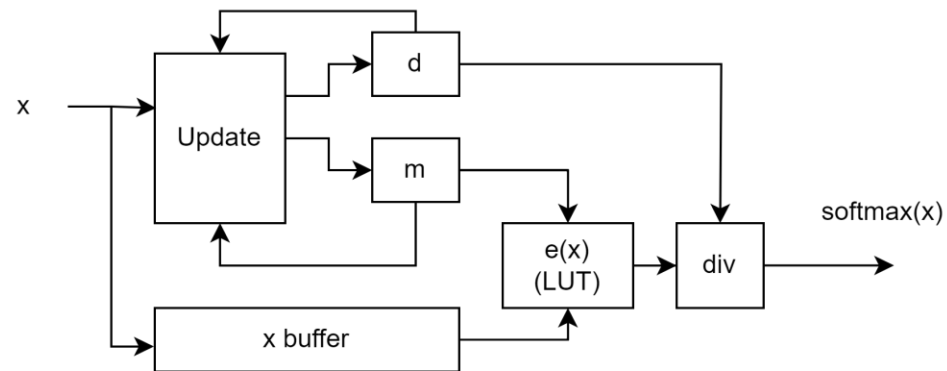
- Softmax
 - 3-pass to 2-pass

$$m_i \leftarrow \max(m_{i-1}, x_i) \quad m_i \leftarrow \max(m_{i-1}, x_i)$$
$$d'_i \leftarrow d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i}$$

$$d_i \leftarrow d_{i-1} + e^{x_i - m_N}$$

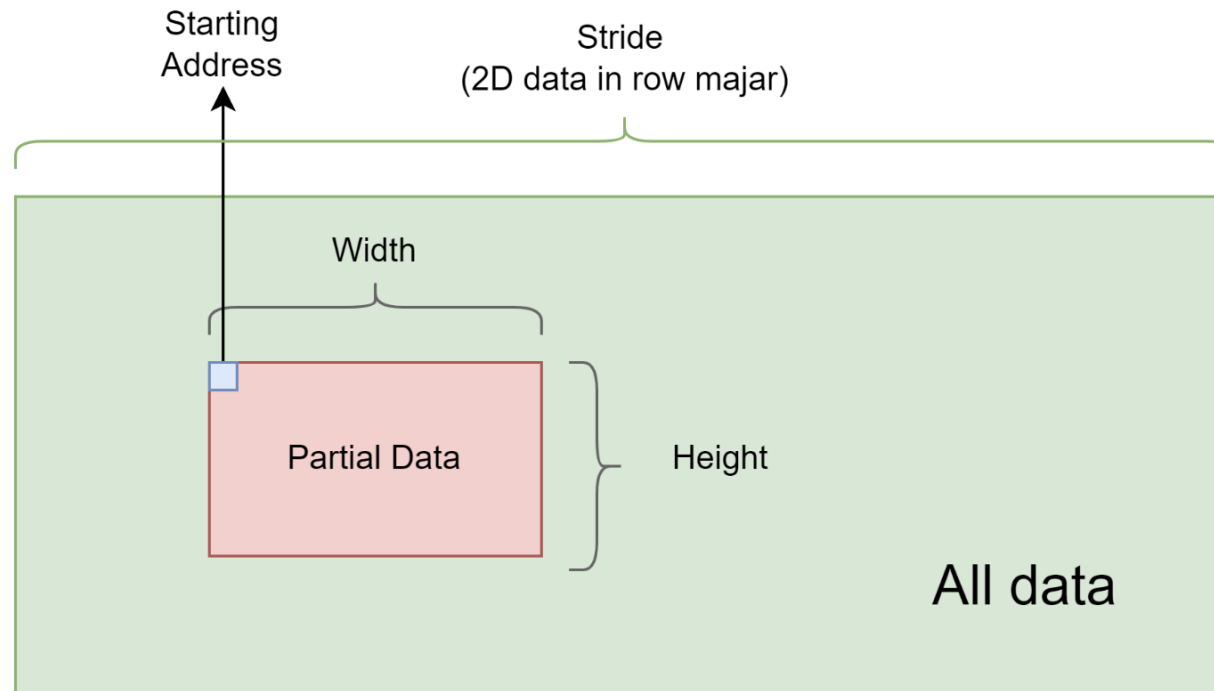
$$a_i \leftarrow \frac{e^{x_i - m_N}}{d_N}$$

$$a_i \leftarrow \frac{e^{x_i - m_N}}{d'_N}$$



Loadable – Data Representation

```
typedef struct Data {  
    uint8_t *addr;           // starting address of the first element  
    uint32_t height;         // number of rows in a channel  
    uint32_t width;          // number of columns in a row  
    uint32_t stride;         // distance between two adjacent rows  
} Data;
```



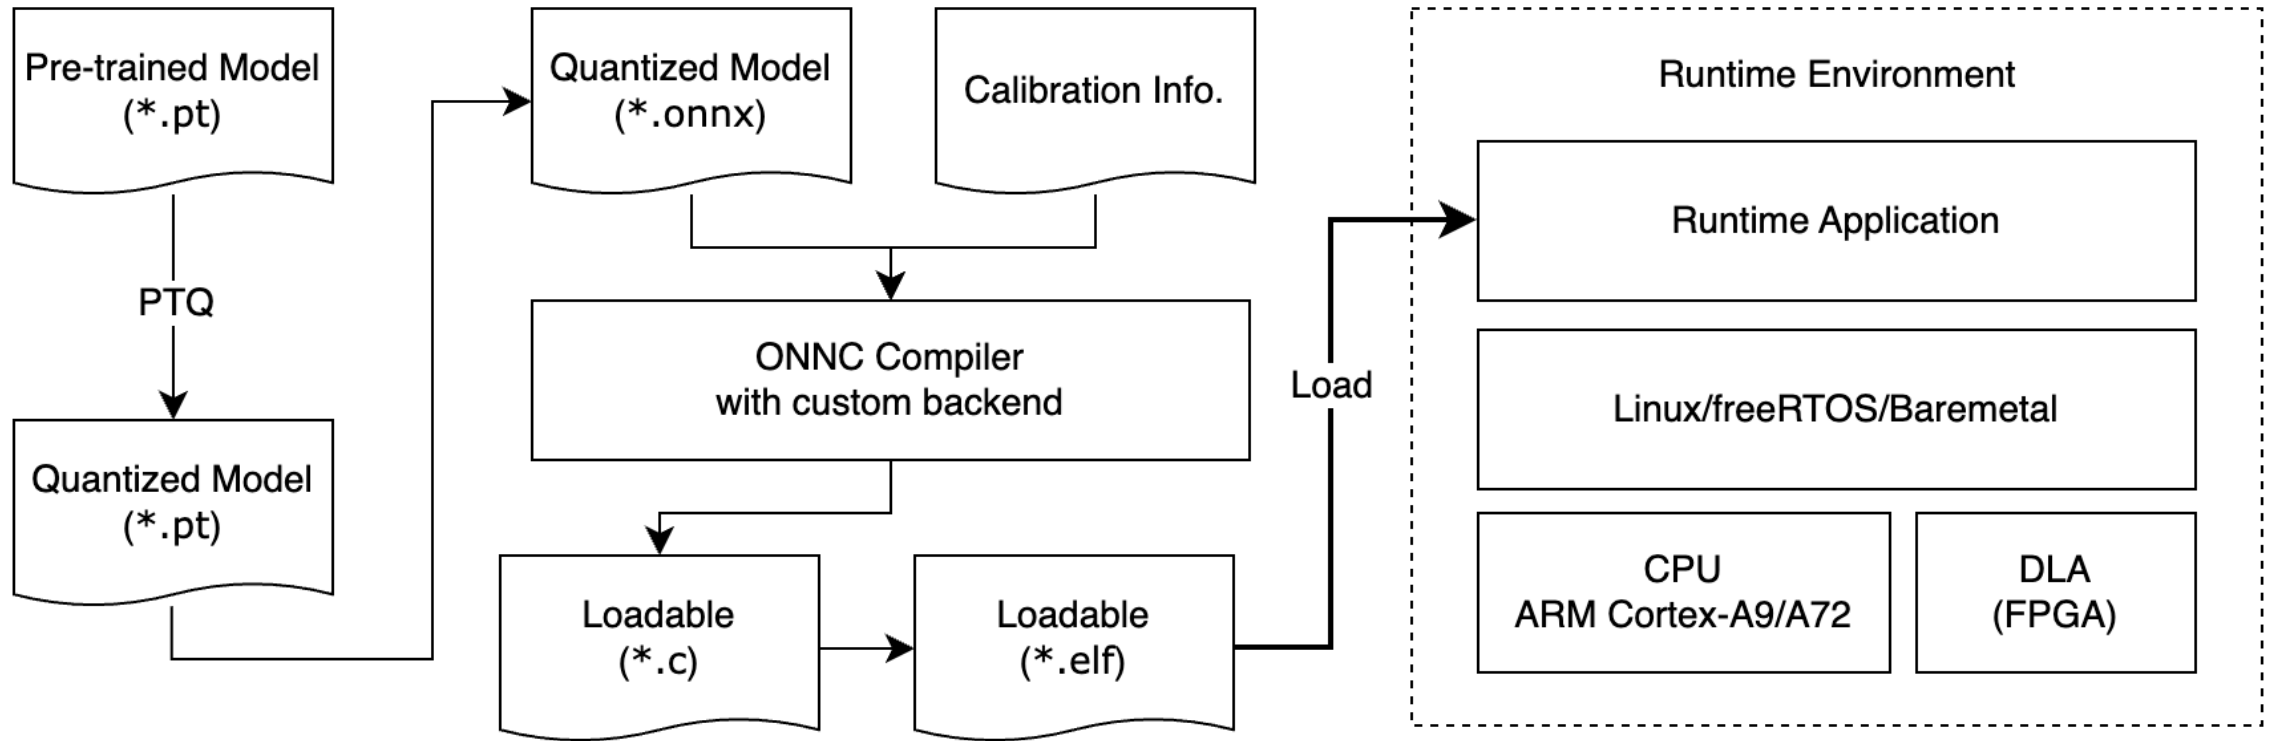
Loadable – Supported Operators

- Conv2D
- GEMM
- MaxPool
- AvgPool
- Softmax
- ReLU
- ReLU6

```
typedef struct ConvArgs {  
    Data *x;           // input  
    Data *W;           // weight  
    Data *B;           // bias  
    Data *y;           // output  
    uint32_t padding;  
    uint32_t stride;  
} ConvArgs;
```

```
typedef struct Maxpool2DArgs {  
    Data *x;           // input  
    Data *y;           // output  
    uint32_t kern_shape; // assuming square kernel (only supports 2 or 3)  
    uint32_t padding;  
    uint32_t stride;  
} MaxpoolArgs;
```

Compilation Flow



Quantization

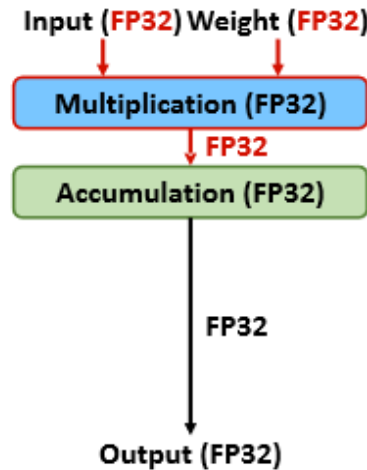
- Quantization method
 - Post training quantization
 - Uniform quantization
 - Symmetric quantization
- Granularity
 - Per-channel quantization
- Integer-only arithmetic

[1] [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)

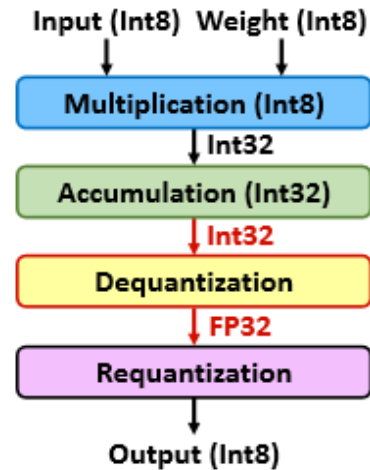
[2] [HAWQV3: Dyadic Neural Network Quantization](#)

[3] [Trained Quantization Thresholds for Accurate and Efficient Fixed-Point Inference of Deep Neural Networks](#)

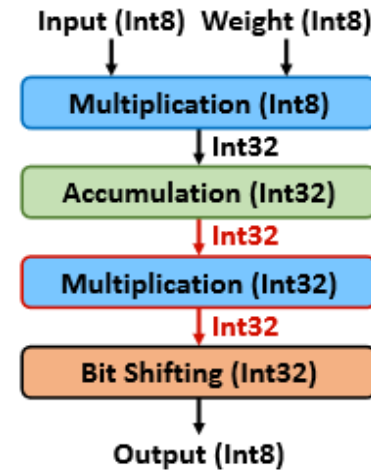
Quantization



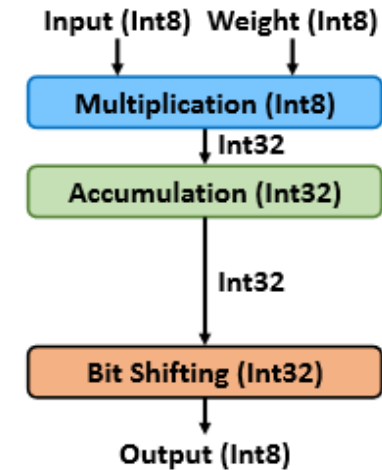
(a) Full-precision.



(b) Simulated quant.



(c) Integer-only quant.



(d) Fixed-point quant.

- Integer-only quant

Paper [1][2] : fixed-point multiplication & bit shifting

$$2^{-n}M_0 := \frac{S_1S_2}{S_3}$$

- n is a non-negative integer
- M_0 is a fixed-point multiplier in the interval $[0.5, 1)$

- Fixed point quant

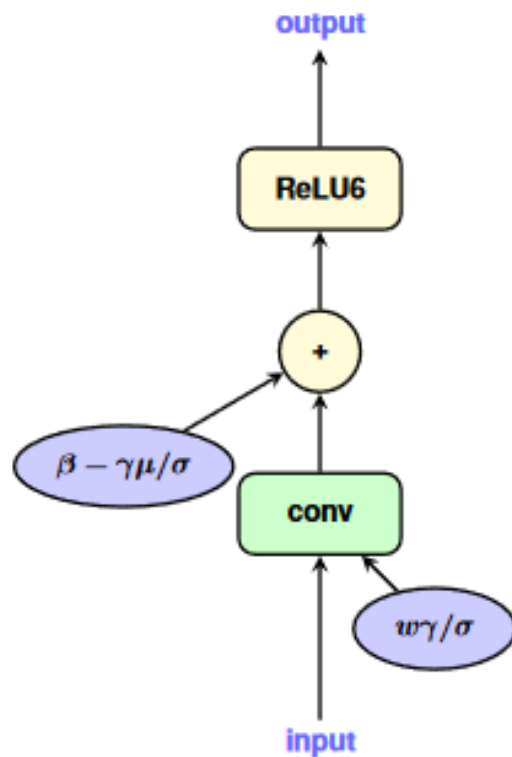
Paper [3] : bit-shift

by constraining scale-factors s_1, s_2, s_3 to strict power-of-2, the scaling operation reduces to a rather simple bit-shift (with round-to-nearest):

$$2^{-f} = \frac{S_1S_2}{S_3}$$

Operator Fusion

Fusing Convolution and Batch Normalization



Batch Normalization

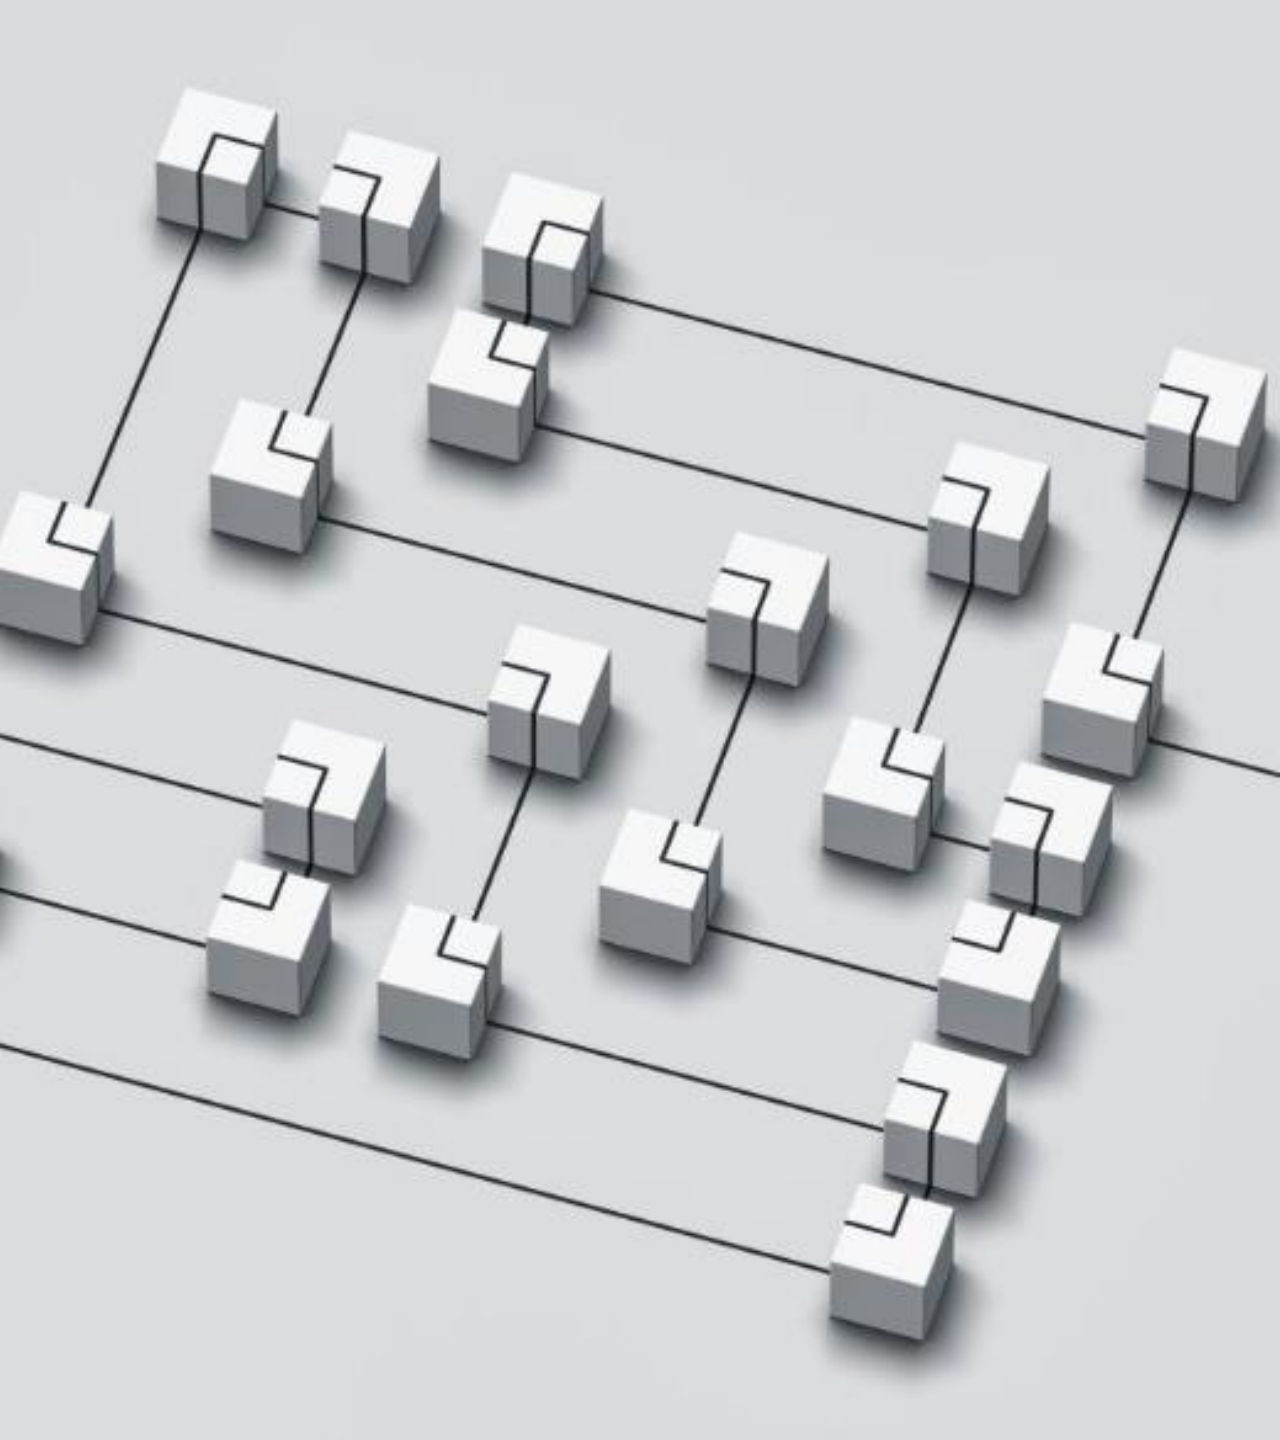
$$y_i = \gamma \left(\frac{x_i - \text{mean}}{\sqrt{\text{variance} + \epsilon}} \right) + \beta$$

Compute effective weights (W_{eff}) by scaling original weights W with the gamma parameter, normalized by the variance plus a small epsilon for stability:

$$W_{\text{eff}} = W \cdot \left(\frac{\gamma}{\sqrt{\text{variance} + \epsilon}} \right)$$

Compute effective bias (b_{eff}) by scaling original bias b with gamma, adjusting for the mean-weight product, and adding beta:

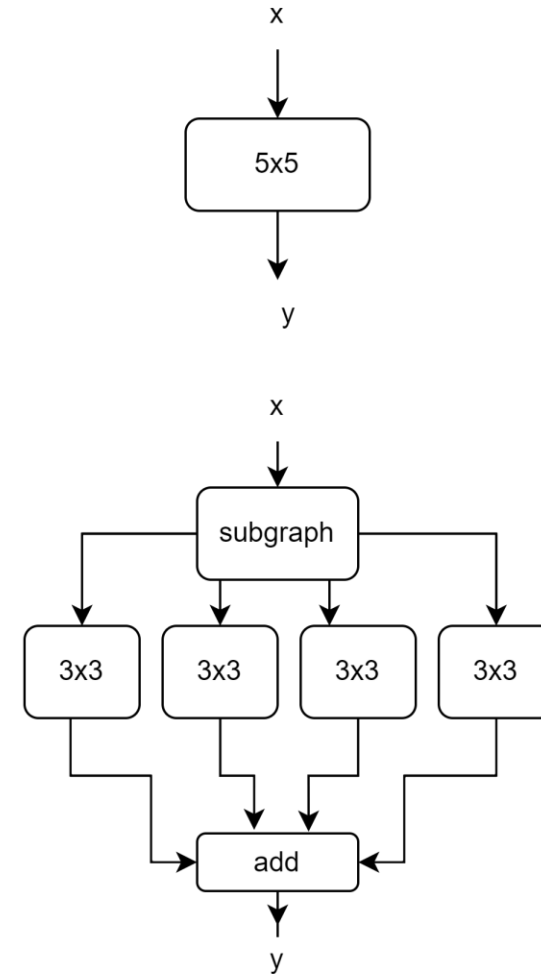
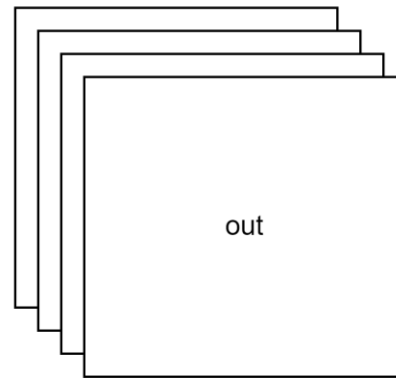
$$b_{\text{eff}} = (b \cdot \gamma) - (\text{mean} \cdot W_{\text{eff}}) + \beta$$



Evaluation

- Inference time (cycle time)
 - Single operator
 - Full model
- Accuracy (ImageNet)
 - Top-1 Accuracy
 - Top-5 Accuracy
- Model size

Question – Software tiling



Question – Memory Requirement

```
1  # N: number of ifmaps/ofmaps
2  + # M: number of filters
3  # H/W: ifmap height/width
4  # R/S: filter height/width
5  # E/F: ofmap height/width
6  # U: stride
7
8  # m: ofmap channels in global buffer
9  # n: number of ifmaps in a pass
10 # e: ofmap width (PE array width)
11 # p: number of filters in a pass
12 # q: (ifmap or filter) channels in a pass
13 # r: number of PE sets for different (ifmap/filter) channels
14 # t: number of PE sets for different filters
15
```

```
[Memory Requirement]
[Name] Alexnet.conv3
=====
[glb_ifmap]          7.0 KiB
[glb_filter]         4.5 KiB
[glb_ofmap]         84.5 KiB
[glb_total]          96.0 KiB
```

```
[Memory Requirement]
[Name] Alexnet.conv2
=====
[glb_ifmap]          3.8 KiB
[glb_filter]         3.1 KiB
[glb_ofmap]         91.1 KiB
[glb_total]          98.0 KiB
```

```
param = dict(
    m = 64,
    n = 1,
    e = 14, # 224
    p = 16, # psum Spad (filter in PE)
    q = 4, # ifmap/filter channels/pass
    r = 1, # different
    t = 4, # (different filter in PE sets)
    W = 56,
    H = 56,
    R = 3,
    S = 3,
    E = 56,
    F = 56,
    C = 64,
    M = 192,
)
m = Mapping(**param)

print(conv2)
print(conv3)
print(m)
```

```
[Memory Requirement]
[Name] mapping_2
=====
[glb_ifmap]          24.5 KiB
[glb_filter]          1.1 KiB
[glb_ofmap]          98.0 KiB
[glb_total]         123.6 KiB
```